

# Outline



A Brief Primer on Query Optimization ✓

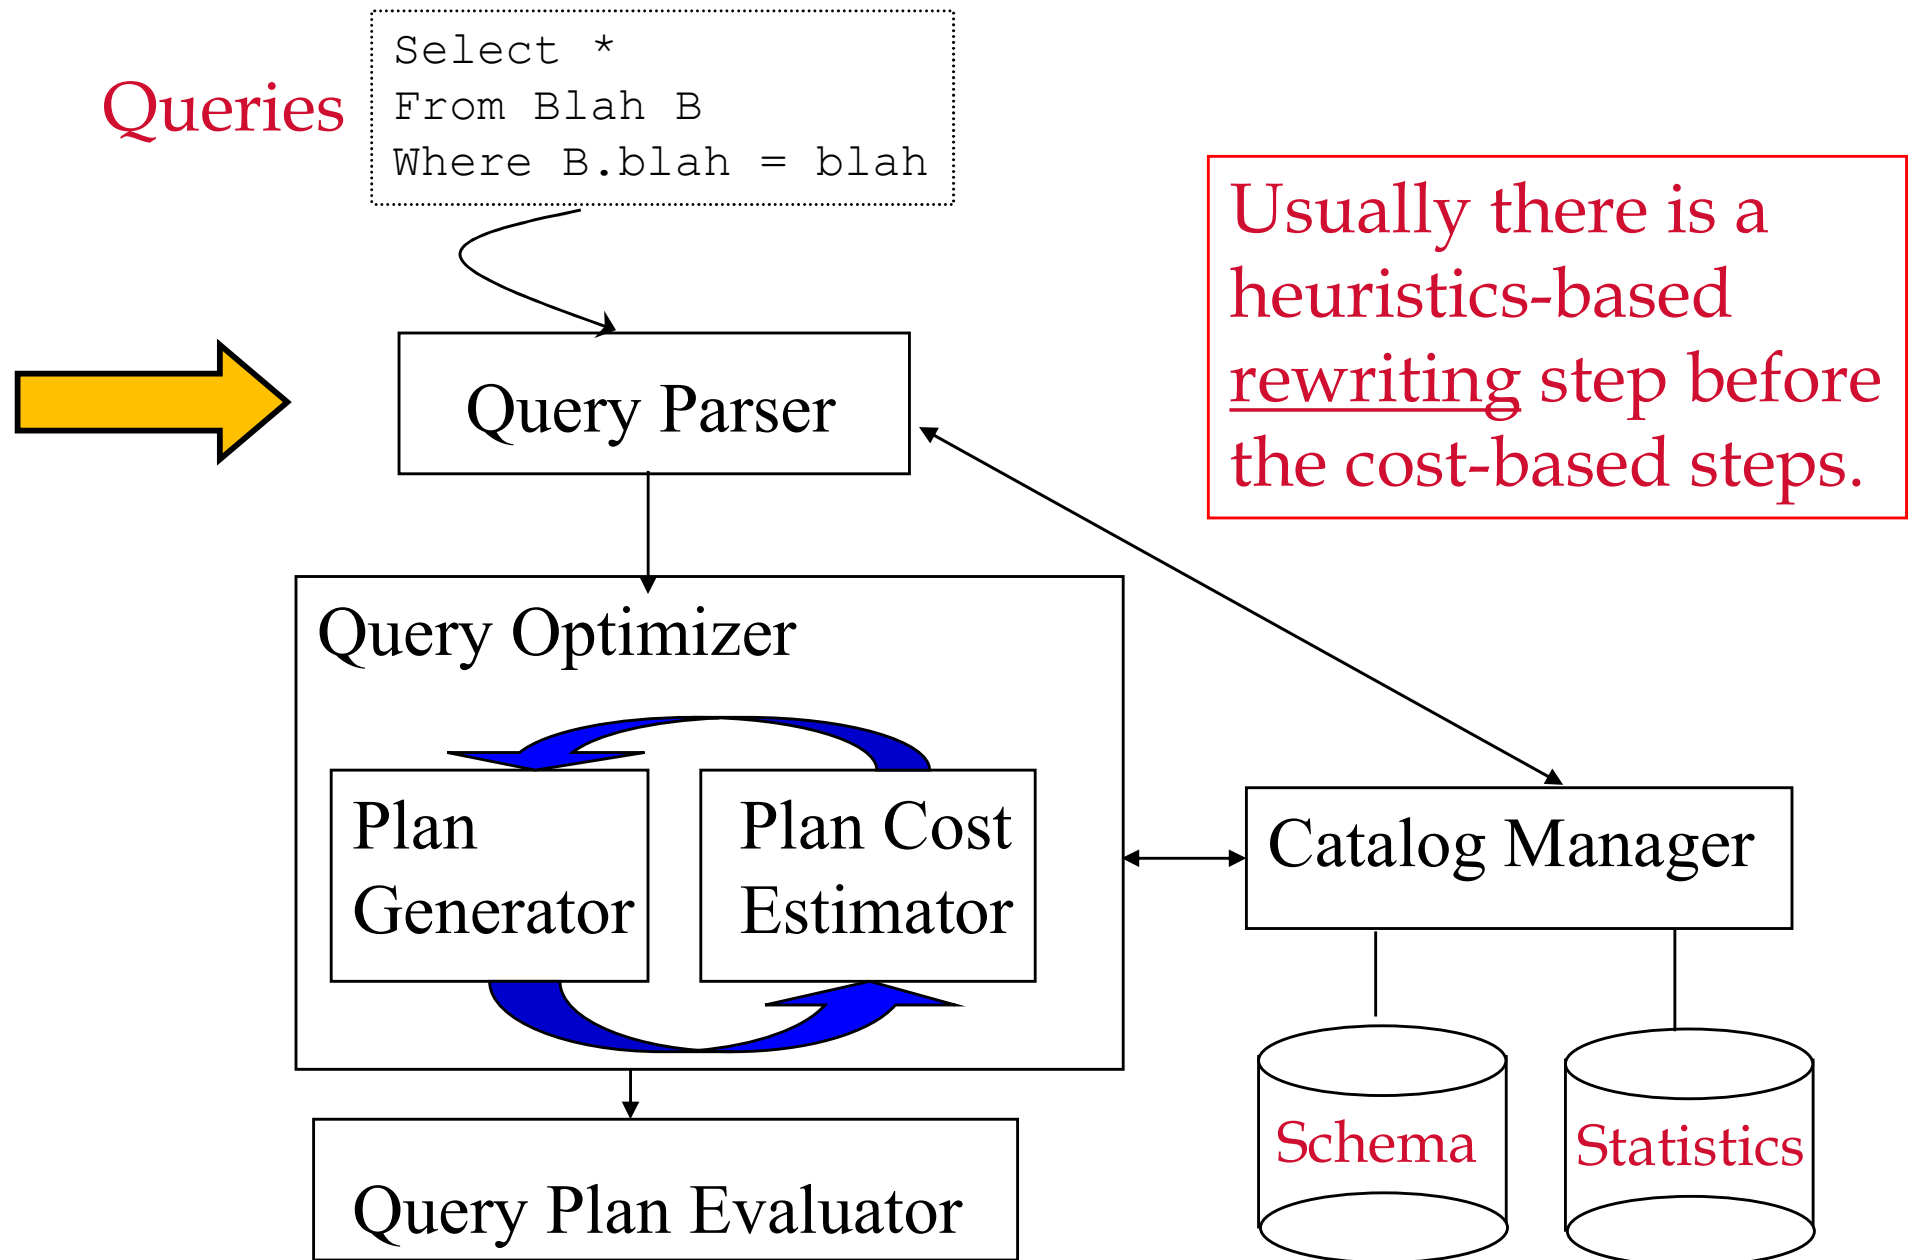
Evaluating Query Plans

Relational Algebra Equivalences

Estimating Plan Costs

Enumerating Plans

# Cost-Based Query Sub-System



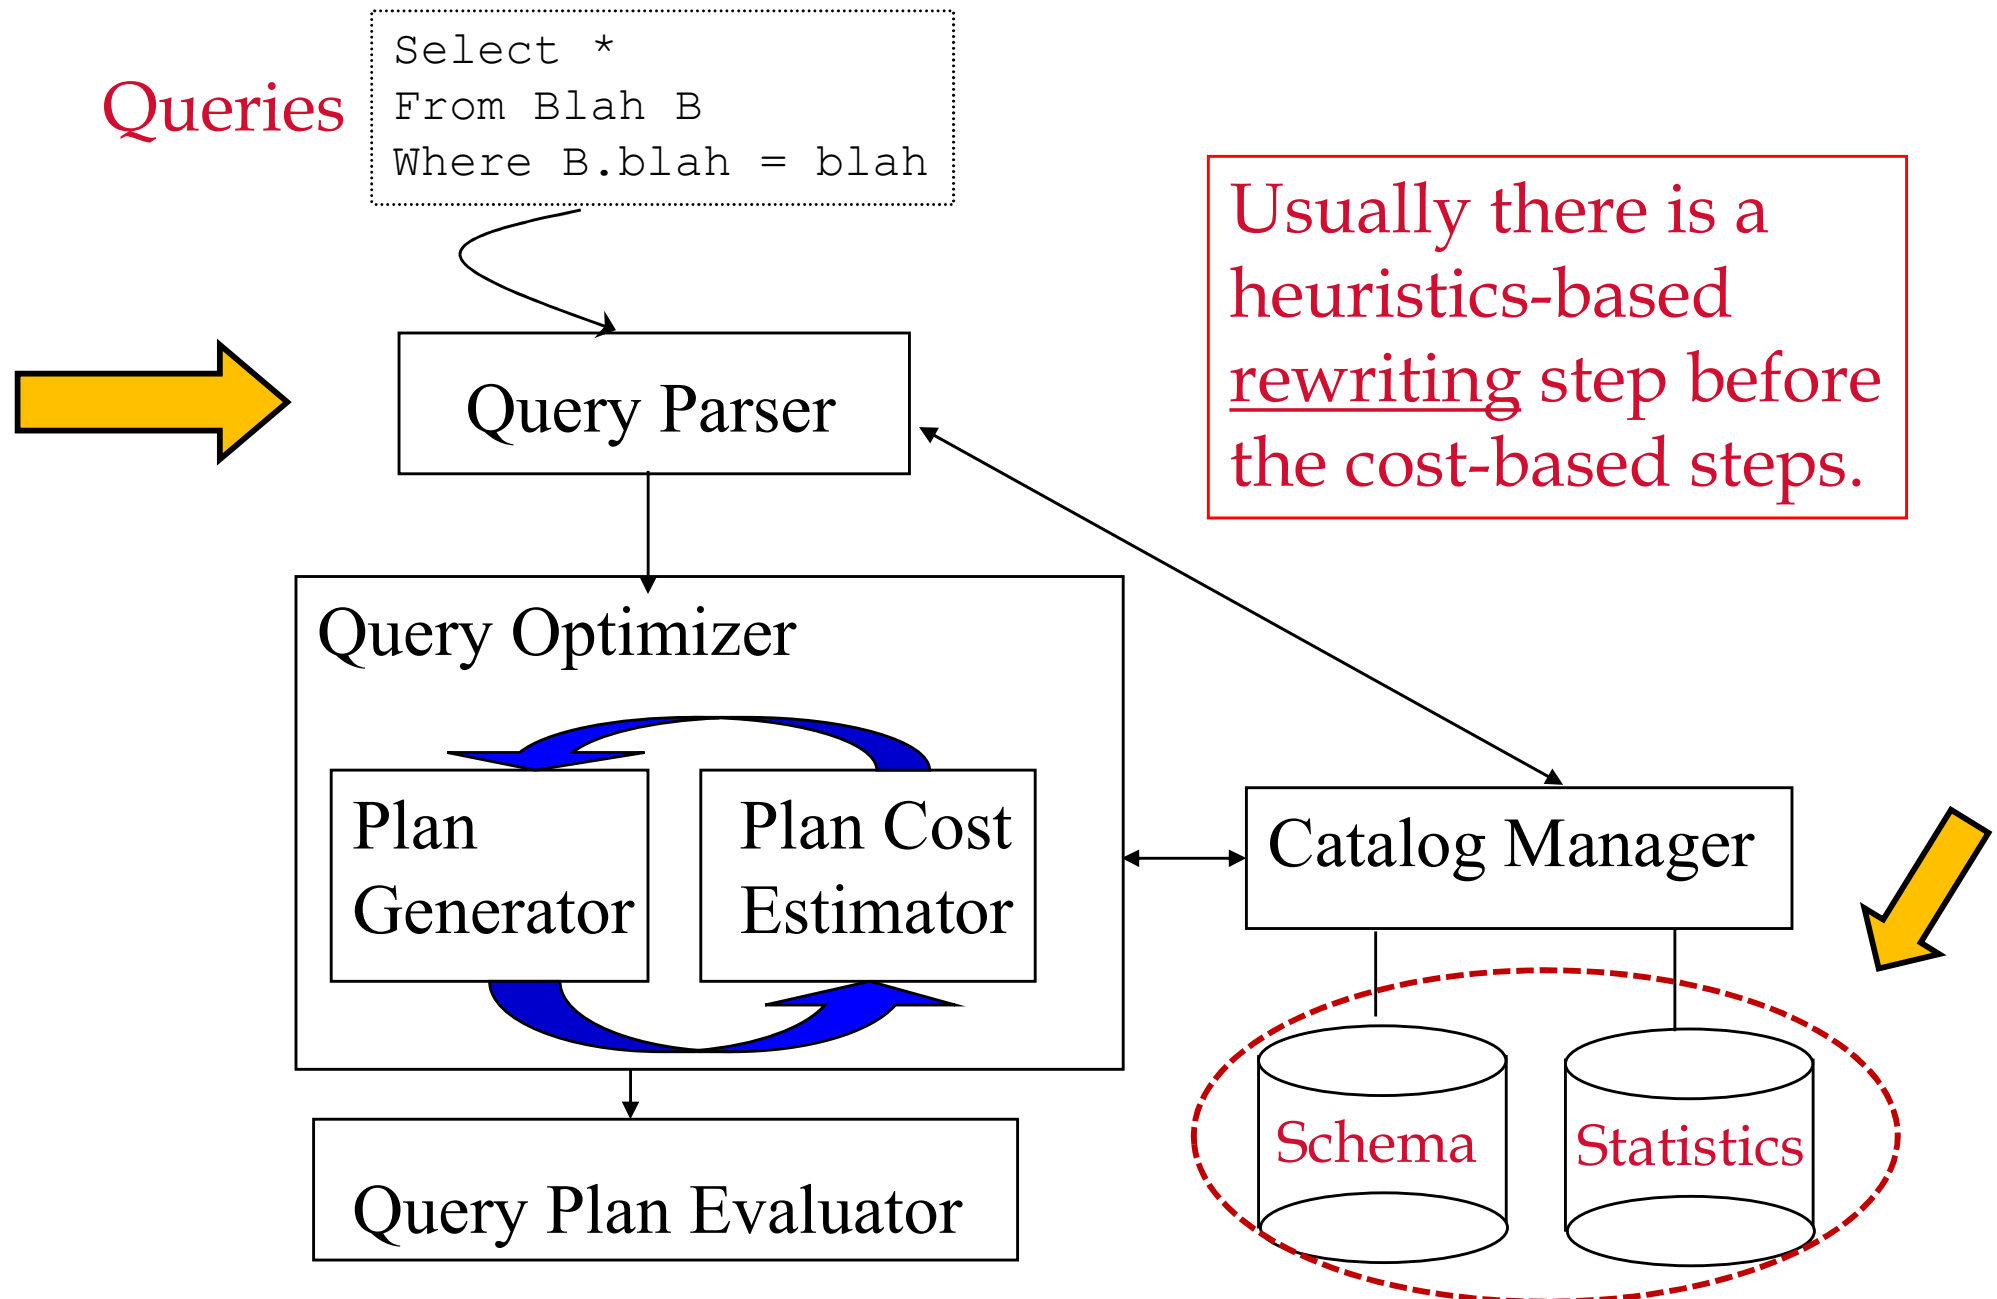
# Query Optimization Steps

- **Step 1:** Queries are parsed into internal forms (e.g., parse trees)
- **Step 2:** Internal forms are transformed into 'canonical forms' (syntactic query optimization)
- **Step 3:** A subset of alternative plans are enumerated
- **Step 4:** Costs for alternative plans are estimated
- **Step 5:** The query evaluation plan with the least estimated cost is picked

# Required Information to Evaluate Queries

- To estimate the costs of query plans, the query optimizer examines the system catalog and retrieves:
  - Information about the types and lengths of fields
  - Statistics about the referenced relations
  - Access paths (indexes) available for relations
- In particular, the *Schema* and *Statistics* components in the Catalog Manager are inspected to find a good enough query evaluation plan

# Cost-Based Query Sub-System



# Catalog Manager: The Schema

- What kind of information do we store at the Schema?
  - Information about **tables** (e.g., table names and integrity constraints) and **attributes** (e.g., attribute names and types)
  - Information about **indices** (e.g., index structures)
  - Information about **users**
- Where do we store such information?
  - In tables, hence, can be queried like any other tables
  - For example: Attribute\_Cat (attr\_name: **string**, rel\_name: **string**; type: **string**; position: **integer**)

# Catalog Manager: Statistics

- What would you store at the Statistics component?
  - $NTuples(R)$ : # records for table  $R$
  - $NPages(R)$ : # pages for  $R$
  - $NKeys(I)$ : # distinct key values for index  $I$
  - $INPages(I)$ : # pages for index  $I$
  - $IHeight(I)$ : # levels for  $I$
  - $ILow(I), IHigh(I)$ : range of values for  $I$
  - ...
- Such statistics are important for estimating plan costs and result sizes (*to be discussed shortly!*)


# SQL Blocks

- SQL queries are optimized by *decomposing* them into a collection of smaller units, called **blocks**
- A block is an SQL query with:
  - No nesting
  - Exactly 1 SELECT and 1 FROM clauses
  - At most 1 WHERE, 1 GROUP BY and 1 HAVING clauses
- A typical relational query optimizer concentrates on optimizing a single block at a time



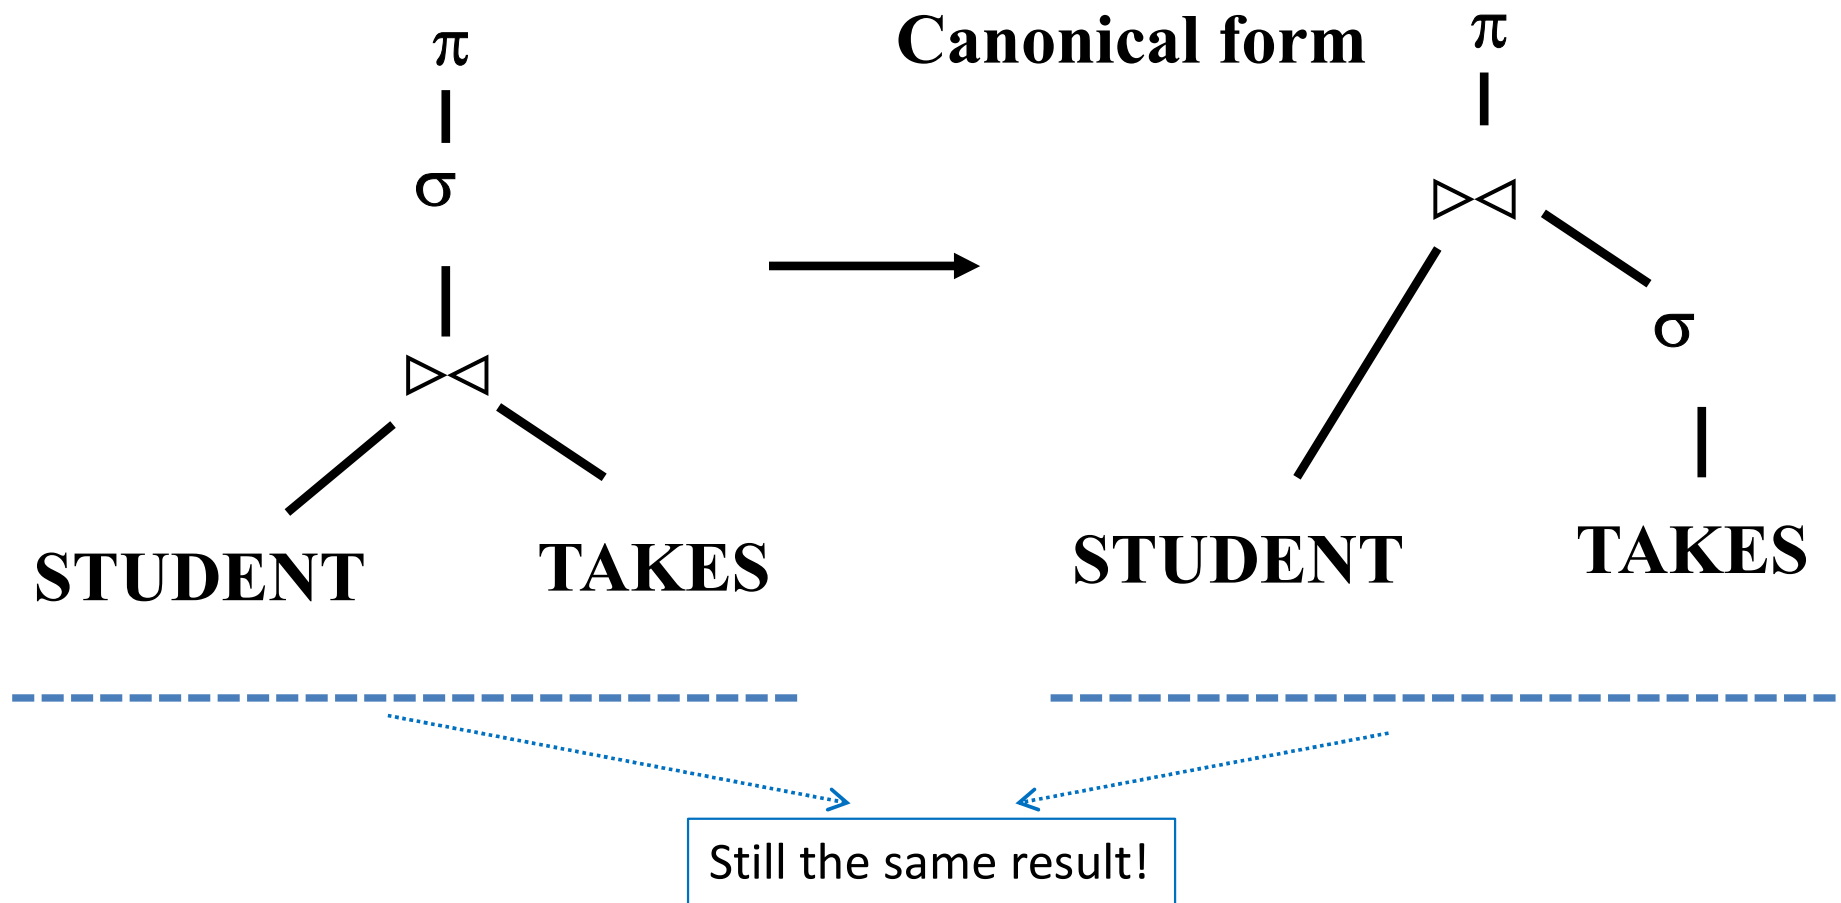
# Translating SQL Queries Into Relational Algebra Trees

**select name  
from STUDENT, TAKES  
where c-id='415' and  
STUDENT.ssn=TAKES.ssn**



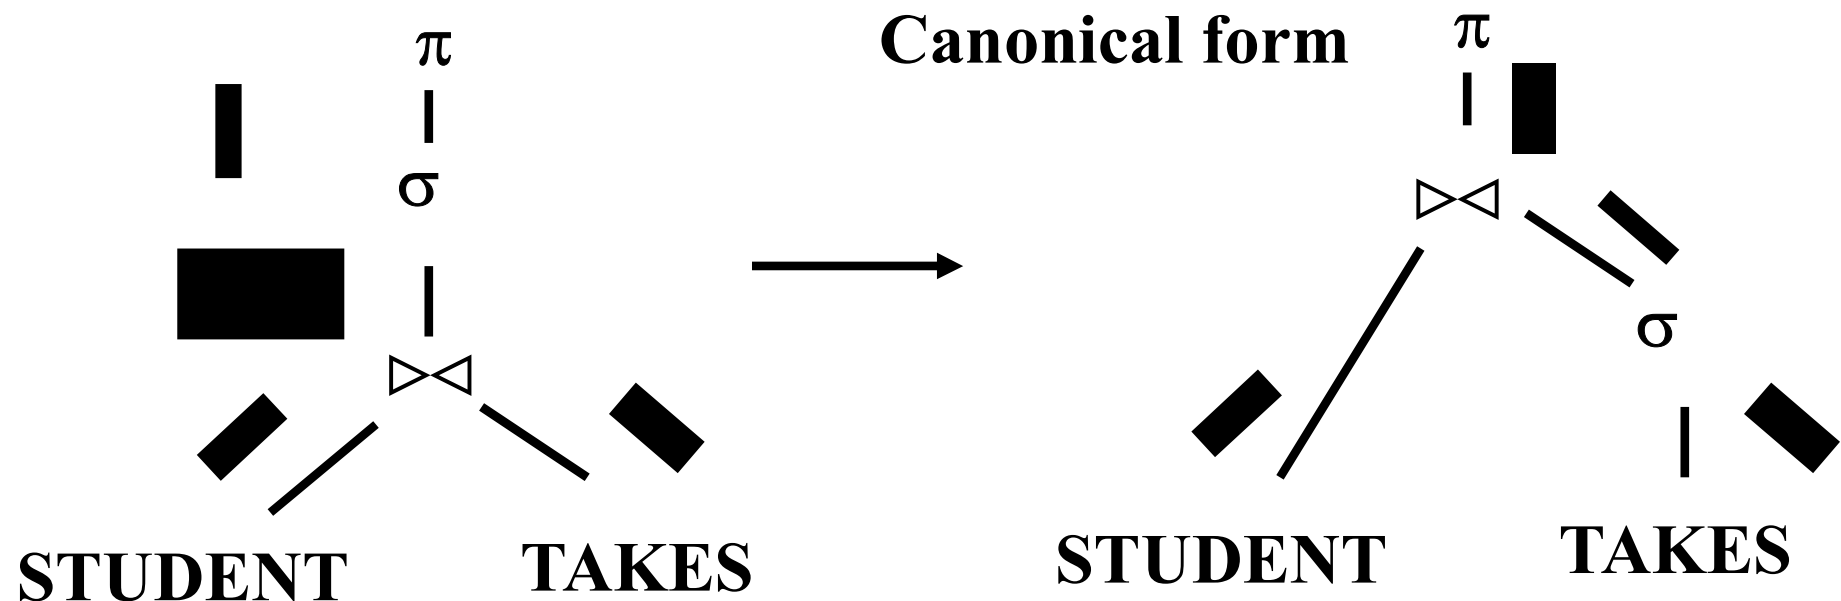
- An SQL block can be thought of as an algebra expression containing:
  - A cross-product of all relations in the FROM clause
  - Selections in the WHERE clause
  - Projections in the SELECT clause
- Remaining operators can be carried out on the result of such SQL block

# Translating SQL Queries Into Relational Algebra Trees (*Cont'd*)




How can this be guaranteed?

# Translating SQL Queries Into Relational Algebra Trees (*Cont'd*)



OBSERVATION: try to perform selections and projections early!

# Translating SQL Queries Into Relational Algebra Trees (*Cont'd*)



How to evaluate a query plan (as opposed to evaluating an operator)?

# Outline



A Brief Primer on Query Optimization

Evaluating Query Plans ✓

Relational Algebra Equivalences

Estimating Plan Costs

Enumerating Plans

# Query Evaluation Plans

- A *query evaluation plan* (or simply a *plan*) consists of an *extended* relational algebra tree (or simply a tree)
- A plan tree consists of annotations at each node indicating:
  - The access methods to use for each relation
  - The implementation method to use for each operator
- Consider the following SQL query **Q**:

```
SELECT S.sname  
FROM Reserves R, Sailors S  
WHERE R.sid=S.sid AND  
      R.bid=100 AND S.rating>5
```

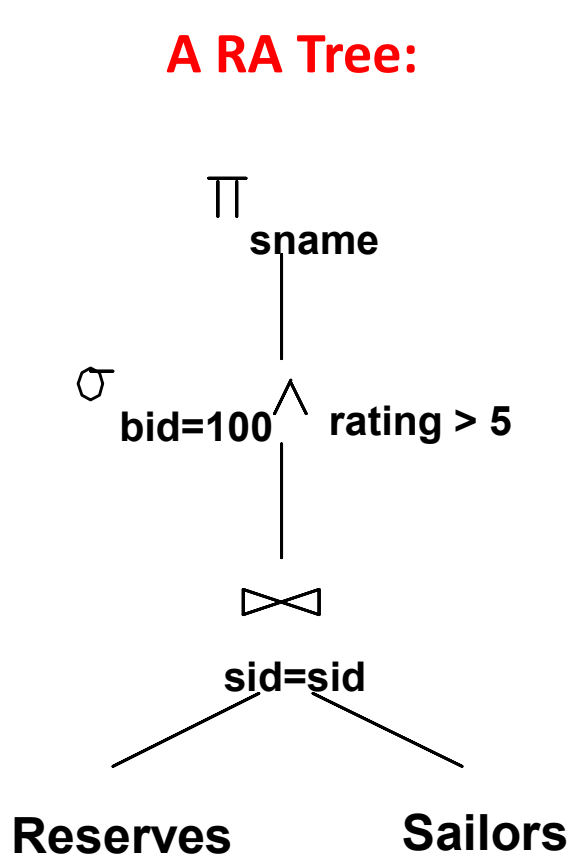
What is the  
corresponding  
RA of **Q**?

# Query Evaluation Plans (Cont'd)

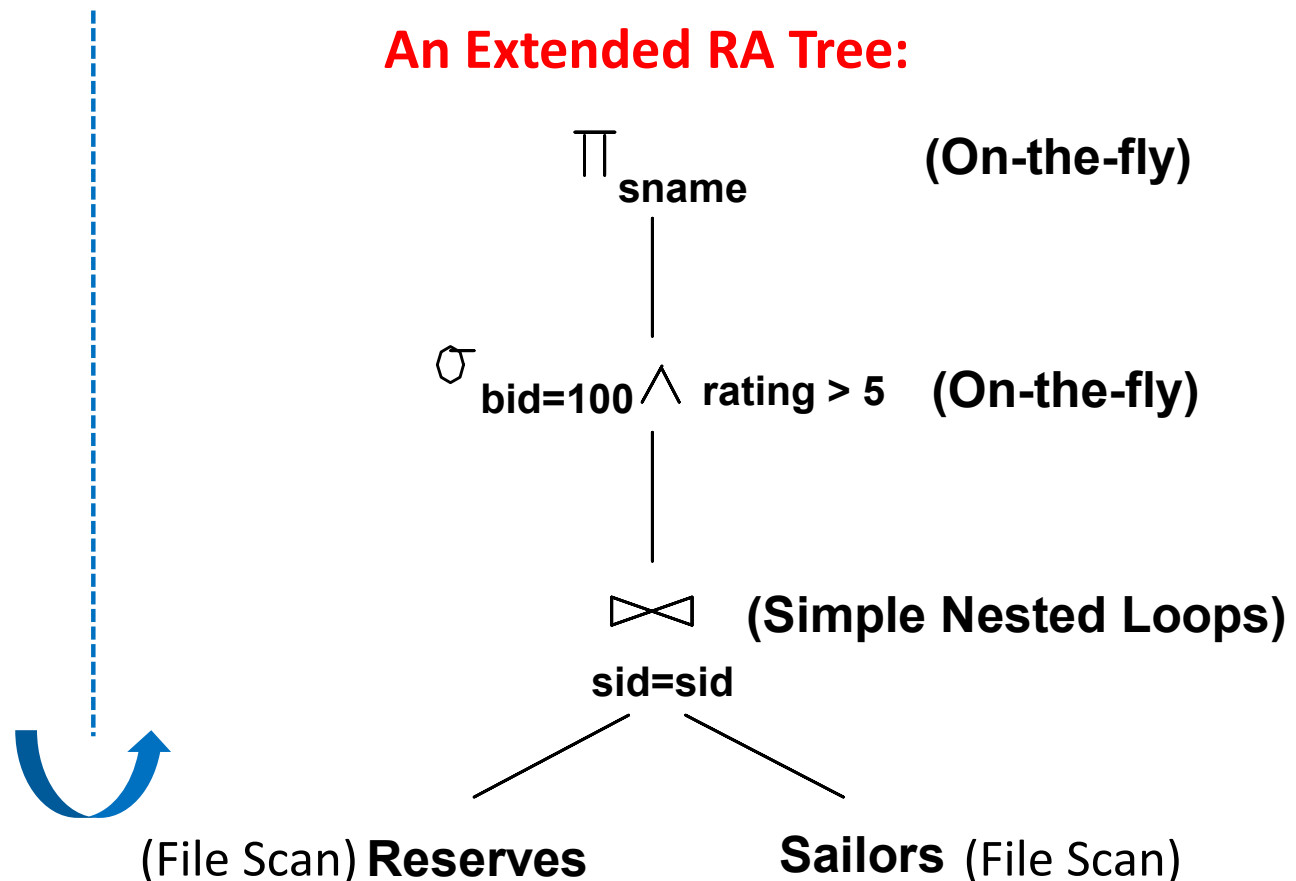
- $Q$  can be expressed in relational algebra as follows:

$$\pi_{sname} (\sigma_{bid=100 \wedge rating > 5} (Reserves \bowtie_{sid=sid} Sailors))$$

**A RA Tree:**

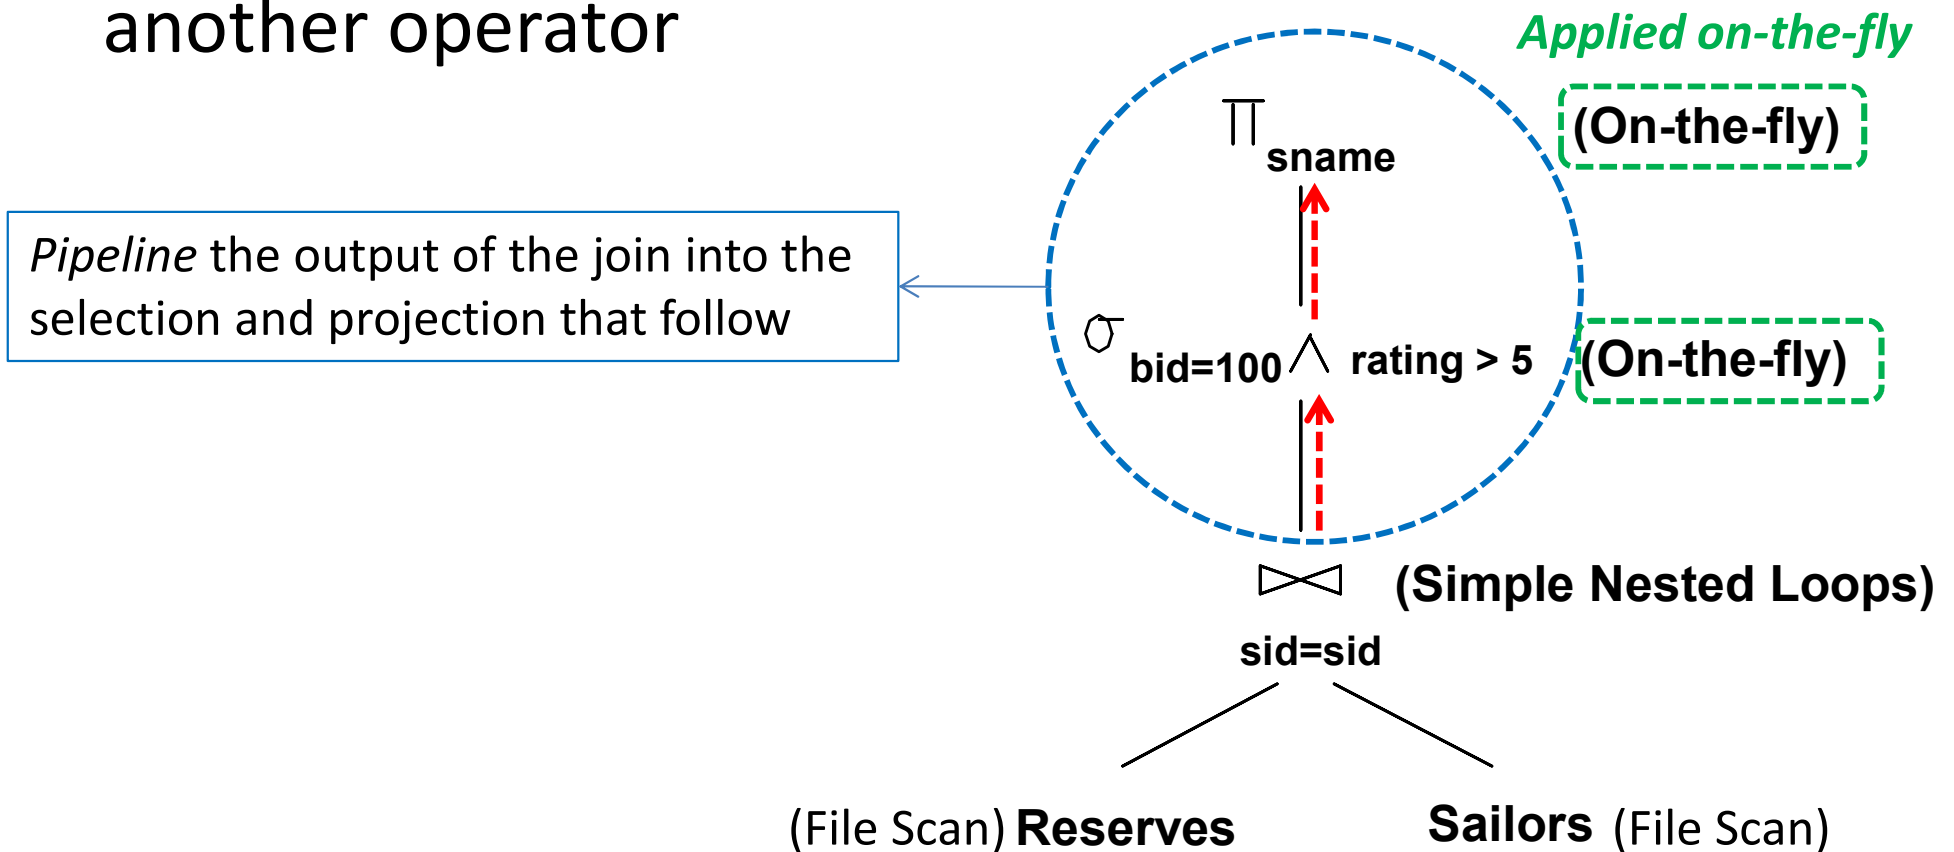


**An Extended RA Tree:**



# Pipelining vs. Materializing

- When a query is composed of several operators, the result of one operator can sometimes be *pipelined* to another operator





# Pipelining vs. Materializing

- When a query is composed of several operators, the result of one operator can sometimes be *pipelined* to another operator

Pipeline the output of the join into the selection and projection that follow

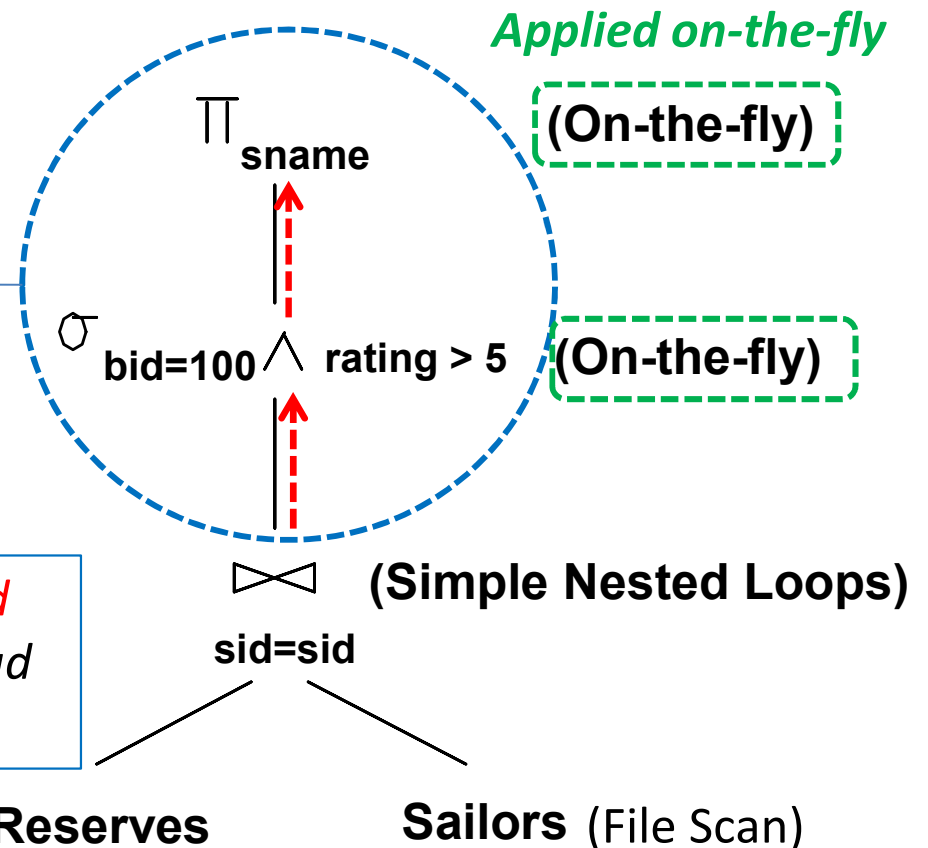


In contrast, a temporary table can be *materialized* to hold the *intermediate result* of the join and *read back* by the selection operation!

(File Scan) **Reserves**

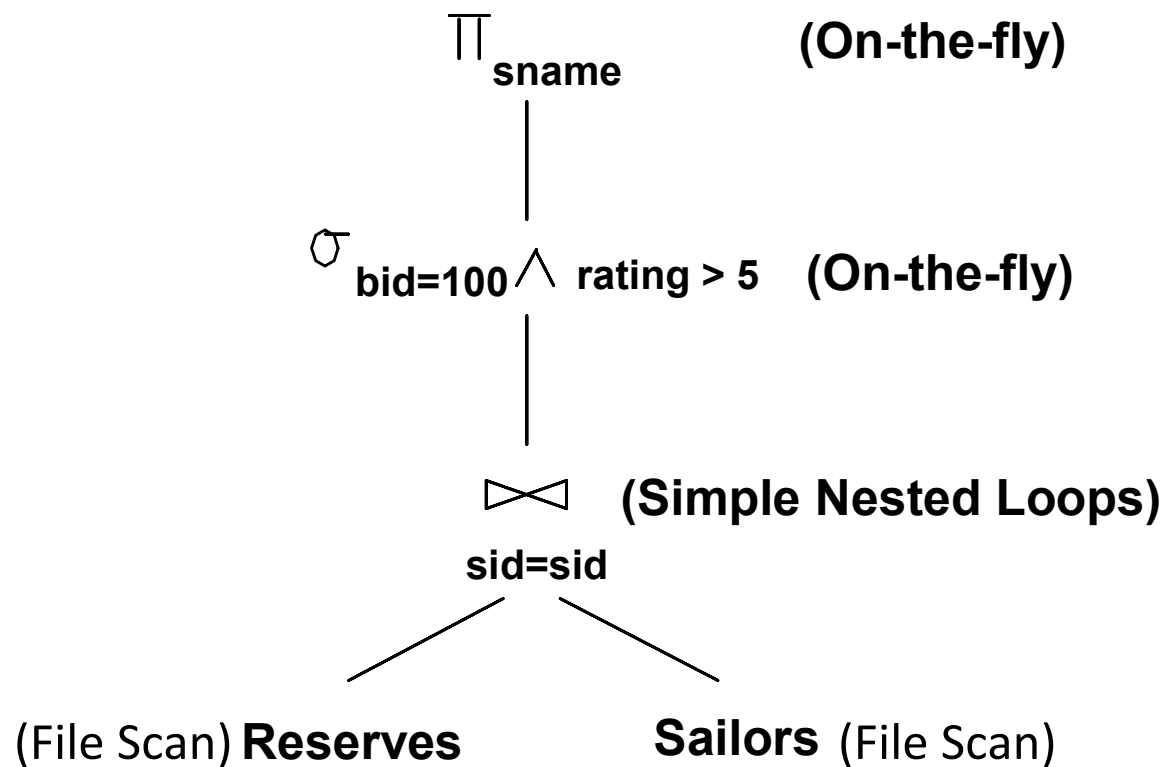
**Sailors** (File Scan)

Pipelining *can* significantly save I/O cost!



# The I/O Cost of the *Q* Plan

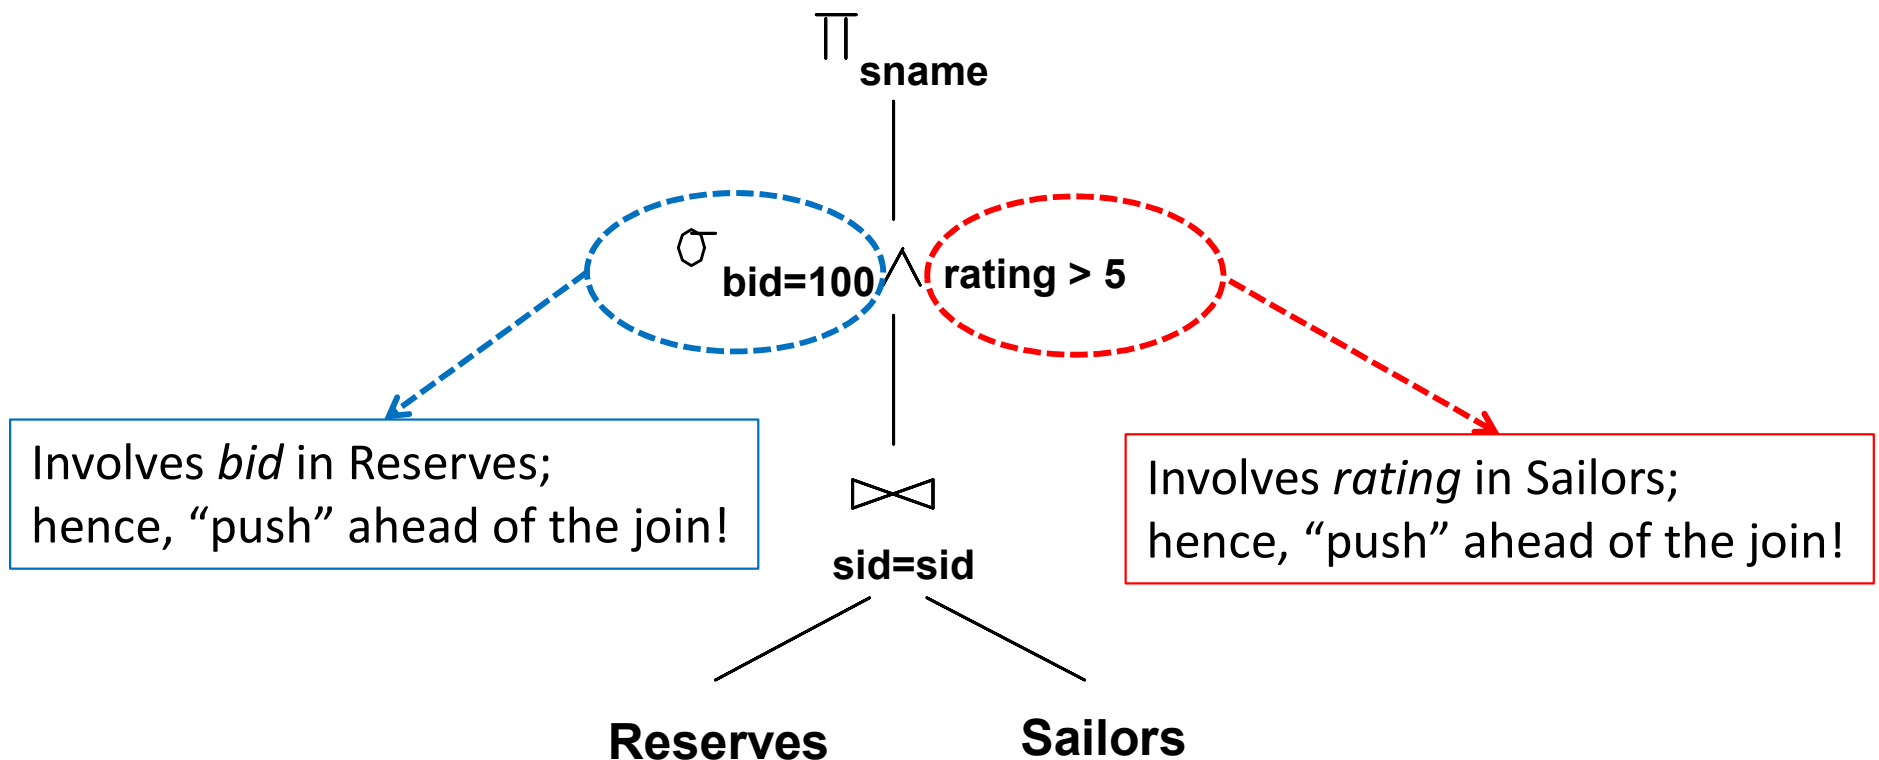
- What is the I/O cost of the following evaluation plan?



- ✓ The cost of the join is  $1000 + 1000 * 500 = 501,000$  I/Os (assuming page-oriented Simple NL join)
- ✓ The selection and projection are done on-the-fly; hence, do not incur additional I/Os

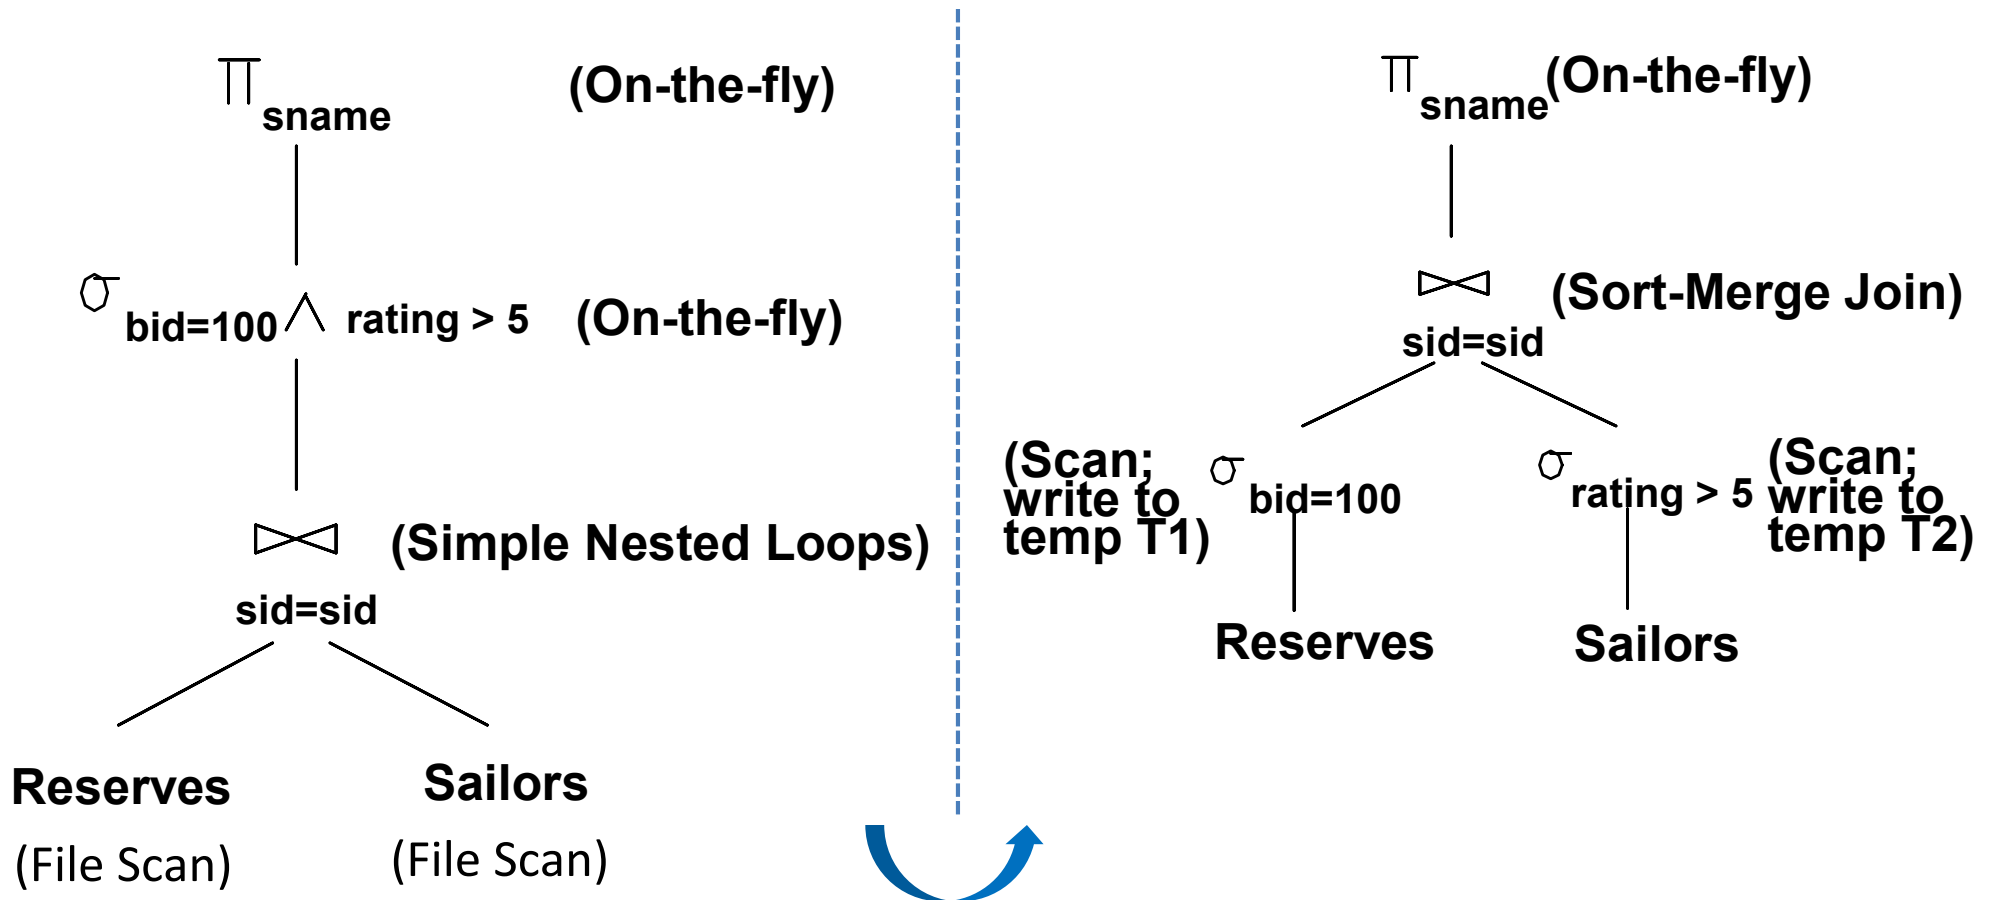
# Pushing Selections

- How can we reduce the cost of a join?
  - By reducing the sizes of the input relations!



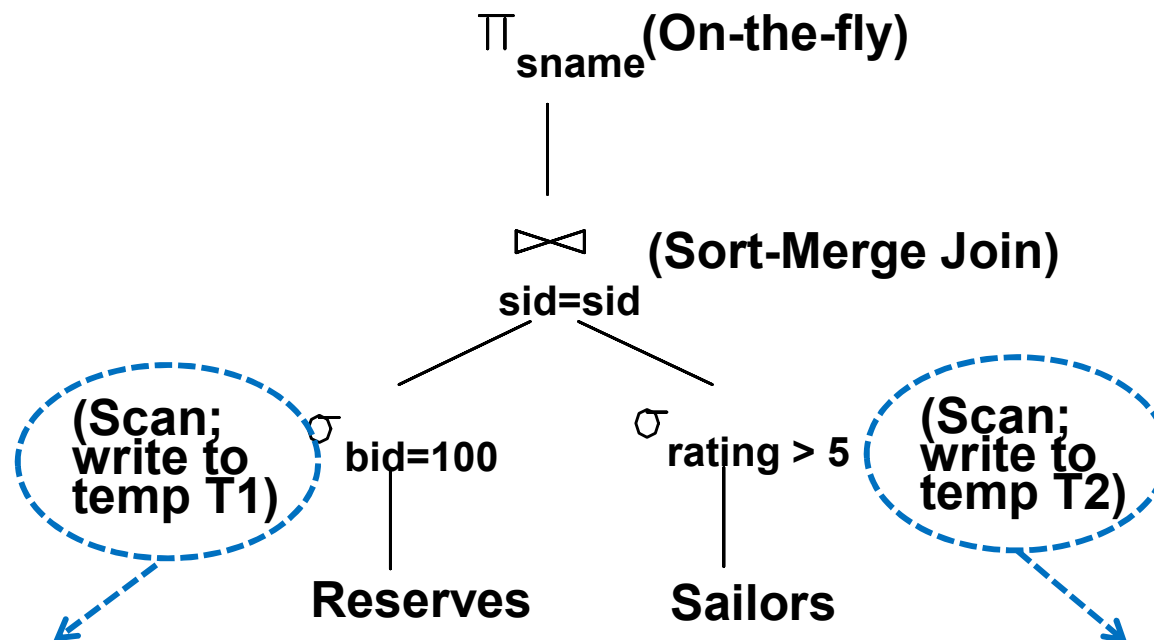
# Pushing Selections

- How can we reduce the cost of a join?
  - By reducing the sizes of the input relations!



# The I/O Cost of the *New Q* Plan

- What is the I/O cost of the following evaluation plan?



**Cost of Scanning Reserves = 1000 I/Os**  
**Cost of Writing T1 = 10\* I/Os (later)**

**Cost of Scanning Sailors = 500 I/Os**  
**Cost of Writing T2 = 250\* I/Os (later)**

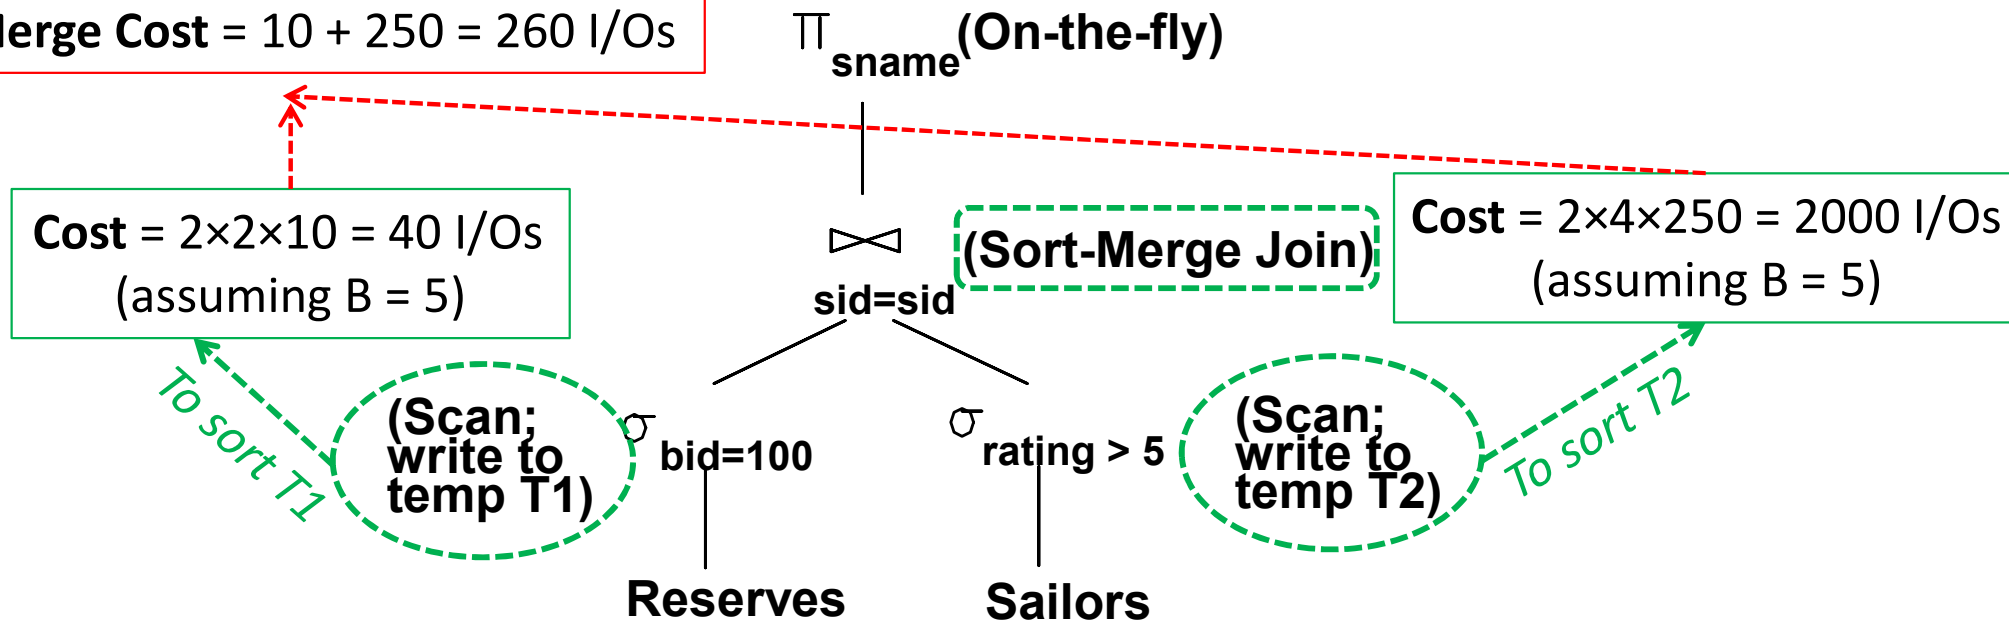
\* Assuming 100 boats and uniform distribution of reservations across boats.

\* Assuming 10 ratings and uniform distribution over ratings.

# The I/O Cost of the *New Q* Plan

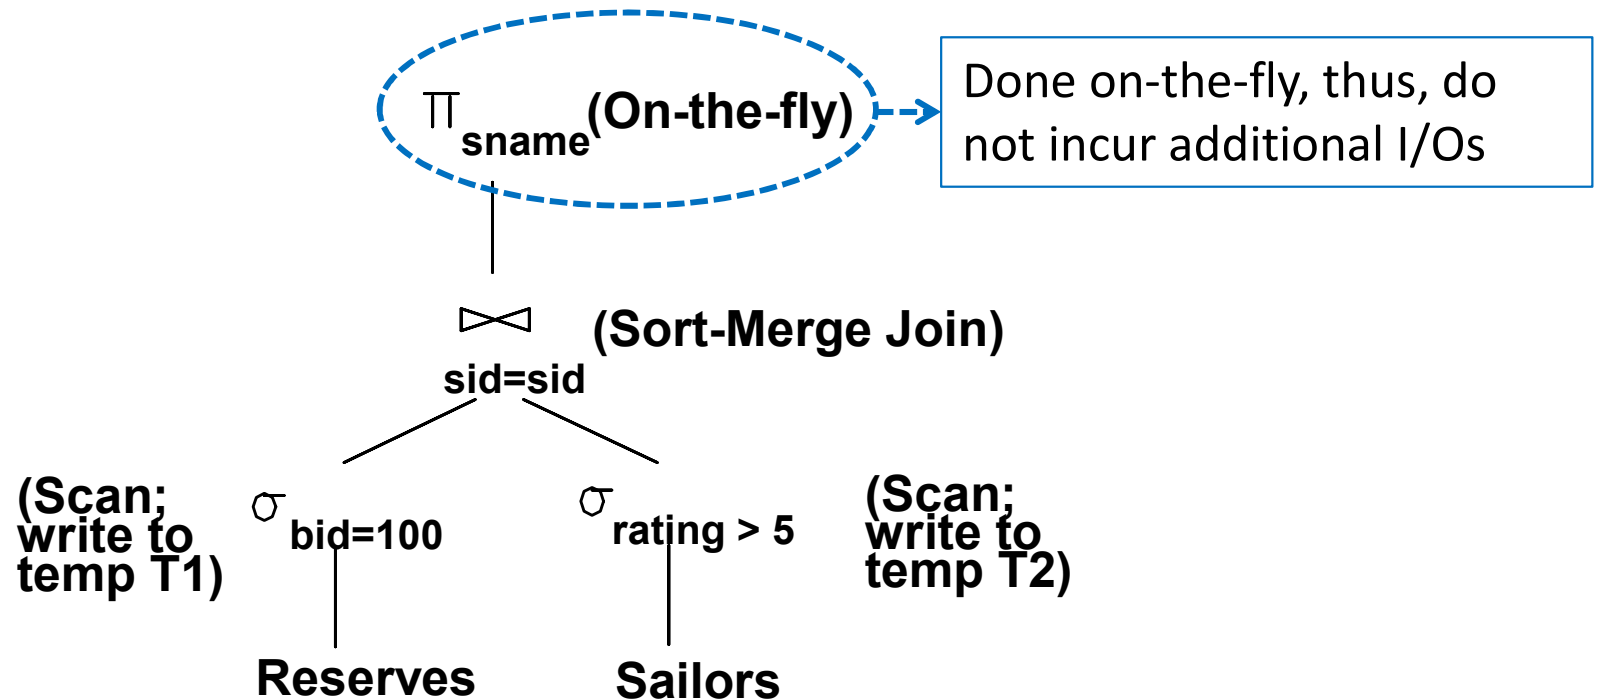
- What is the I/O cost of the following evaluation plan?

Merge Cost =  $10 + 250 = 260$  I/Os



# The I/O Cost of the *New Q* Plan

- What is the I/O cost of the following evaluation plan?



# The I/O Cost of the *New Q* Plan

- What is the I/O cost of the following evaluation plan?

**Merge Cost =  $10 + 250 = 260$  I/Os**

$\Pi_{\text{sname}}$  (On-the-fly)

Done on-the-fly, thus, do not incur additional I/Os

**Cost =  $2 \times 2 \times 10 = 40$  I/Os**  
(assuming  $B = 5$ )

**Cost =  $2 \times 4 \times 250 = 2000$  I/Os**  
(assuming  $B = 5$ )

$\bowtie$  (Sort-Merge Join)  
 $\text{sid}=\text{sid}$

(Scan;  
write to  
temp T1)

$\sigma_{\text{bid}=100}$

Reserves

$\sigma_{\text{rating} > 5}$

Sailors

(Scan;  
write to  
temp T2)

To sort T2

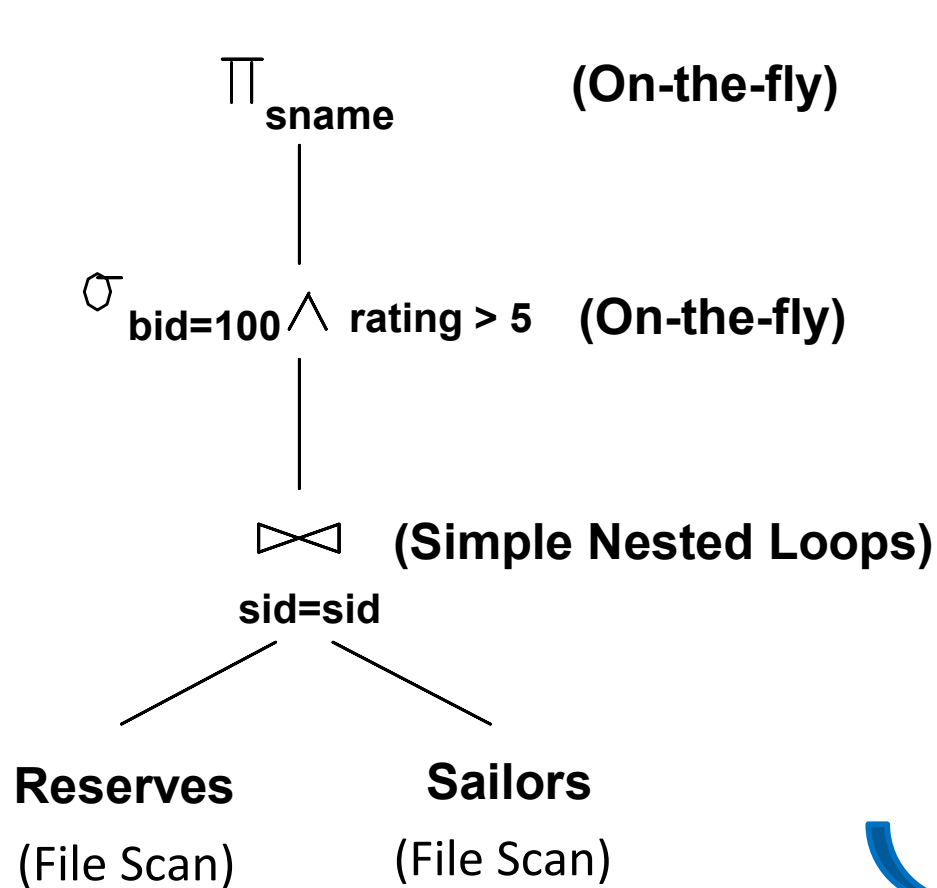
**Cost of Scanning Reserves = 1000 I/Os**  
**Cost of Writing T1 = 10 I/Os (later)**

**Cost of Scanning Sailors = 500 I/Os**  
**Cost of Writing T2 = 250 I/Os (later)**

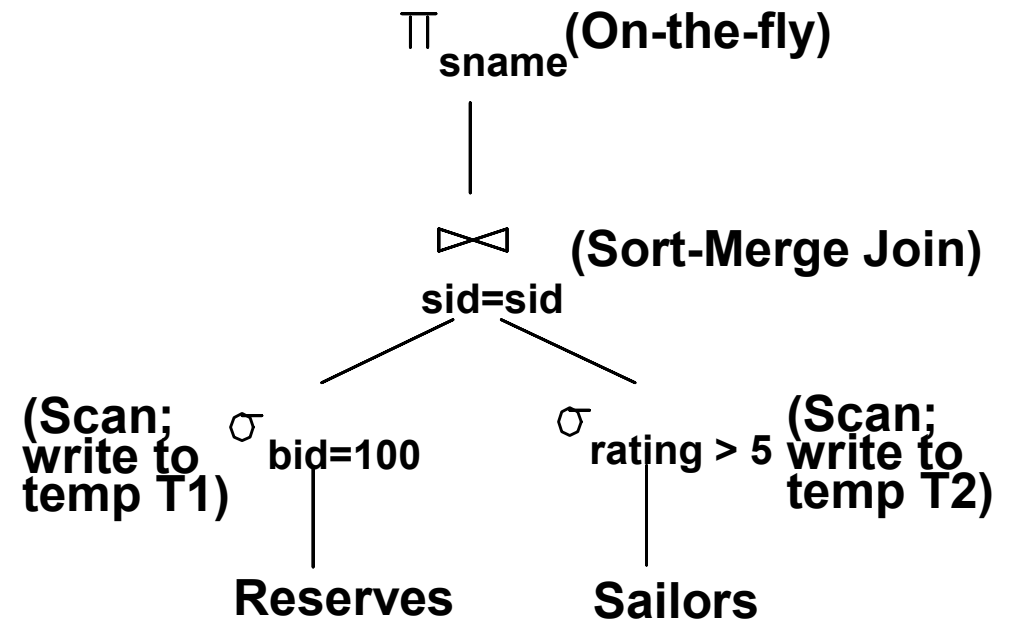
**Total Cost =  $1000 + 10 + 500 + 250 + 40 + 2000 + 260 = 4060$  I/Os**



# The I/O Costs of the *Two Q* Plans




**Total Cost = 501, 000 I/Os**



**Total Cost = 4060 I/Os**

# Pushing Projections

- How can we reduce the cost of a join?
  - By reducing the sizes of the input relations!
- Consider (again) the following plan:




- What are the attributes required from T1 and T2?
  - *Sid* from T1
  - *Sid* and *sname* from T2

Hence, as we scan Reserves and Sailors we can also remove unwanted columns (i.e., “Push” the projections ahead of the join)!

# Pushing Projections


- How can we reduce the cost of a join?
  - By reducing the sizes of the input relations!
- Consider (again) the following plan:



The cost after applying this heuristic can become 2000 I/Os (as opposed to 4060 I/Os with only pushing the selection)!

# Using Indexes

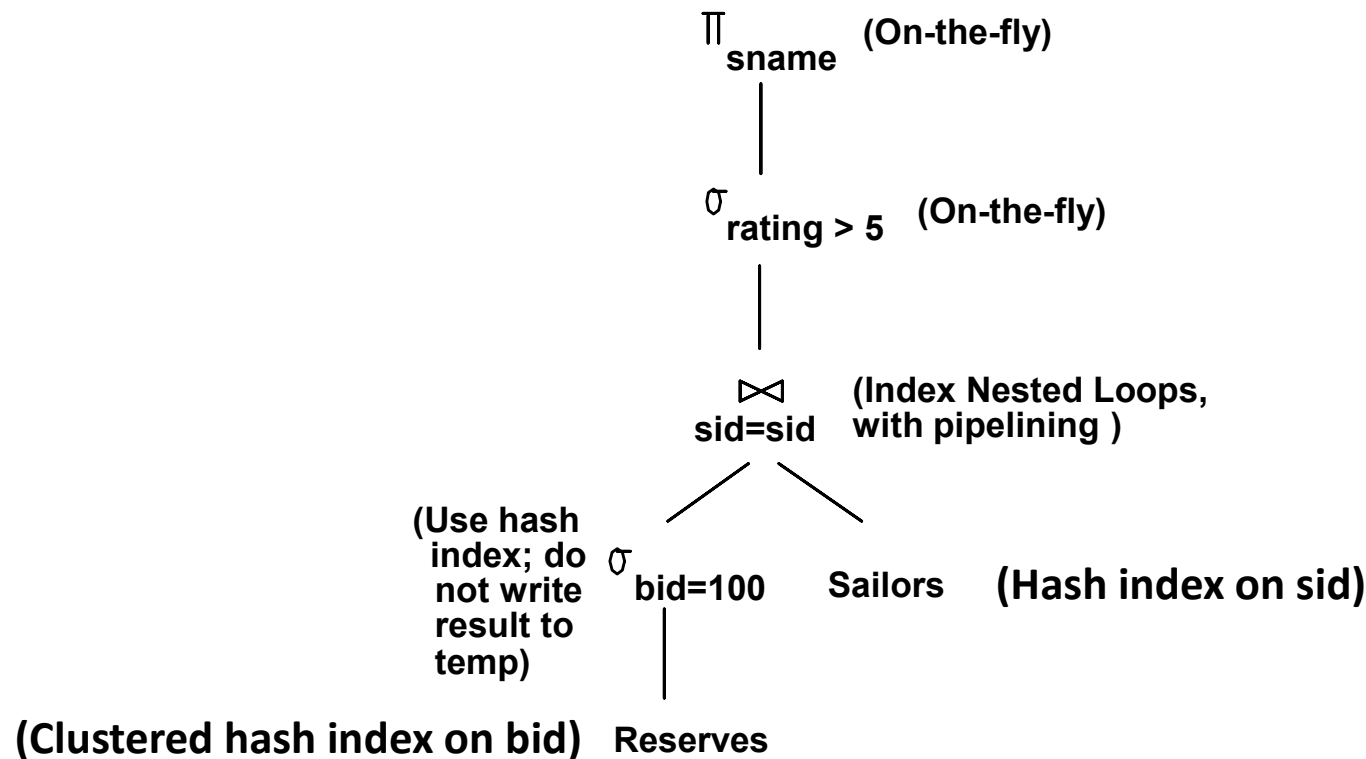
- What if indexes are available on Reserves and Sailors?



- ✓ With clustered index on *bid* of Reserves, we get  $100,000/100 = 1000$  tuples (assuming 100 boats and uniform distribution of reservations across boats)
- ✓ Since the index is clustered, the 1000 tuples appear consecutively within the same bucket; thus # of pages =  $1000/100 = 10$  pages

# Using Indexes

- What if indexes are available on Reserves and Sailors?



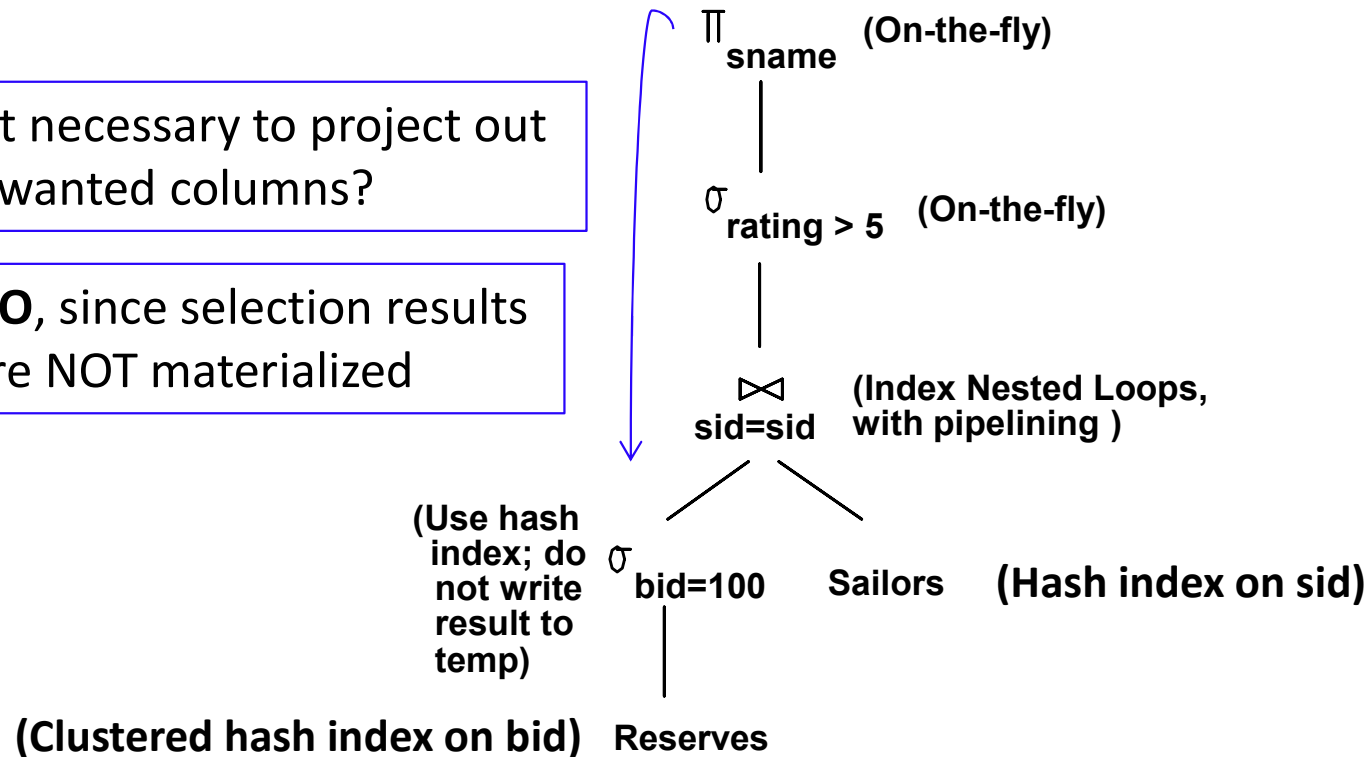
- ✓ For each selected Reserves tuple, we can retrieve matching Sailors tuples using the hash index on the *sid* field
- ✓ Selected Reserves tuples need not be materialized and the join result can be pipelined!
- ✓ For each tuple in the join result, we apply *rating* > 5 and the projection of *sname* on-the-fly

# Using Indexes

- What if indexes are available on Reserves and Sailors?

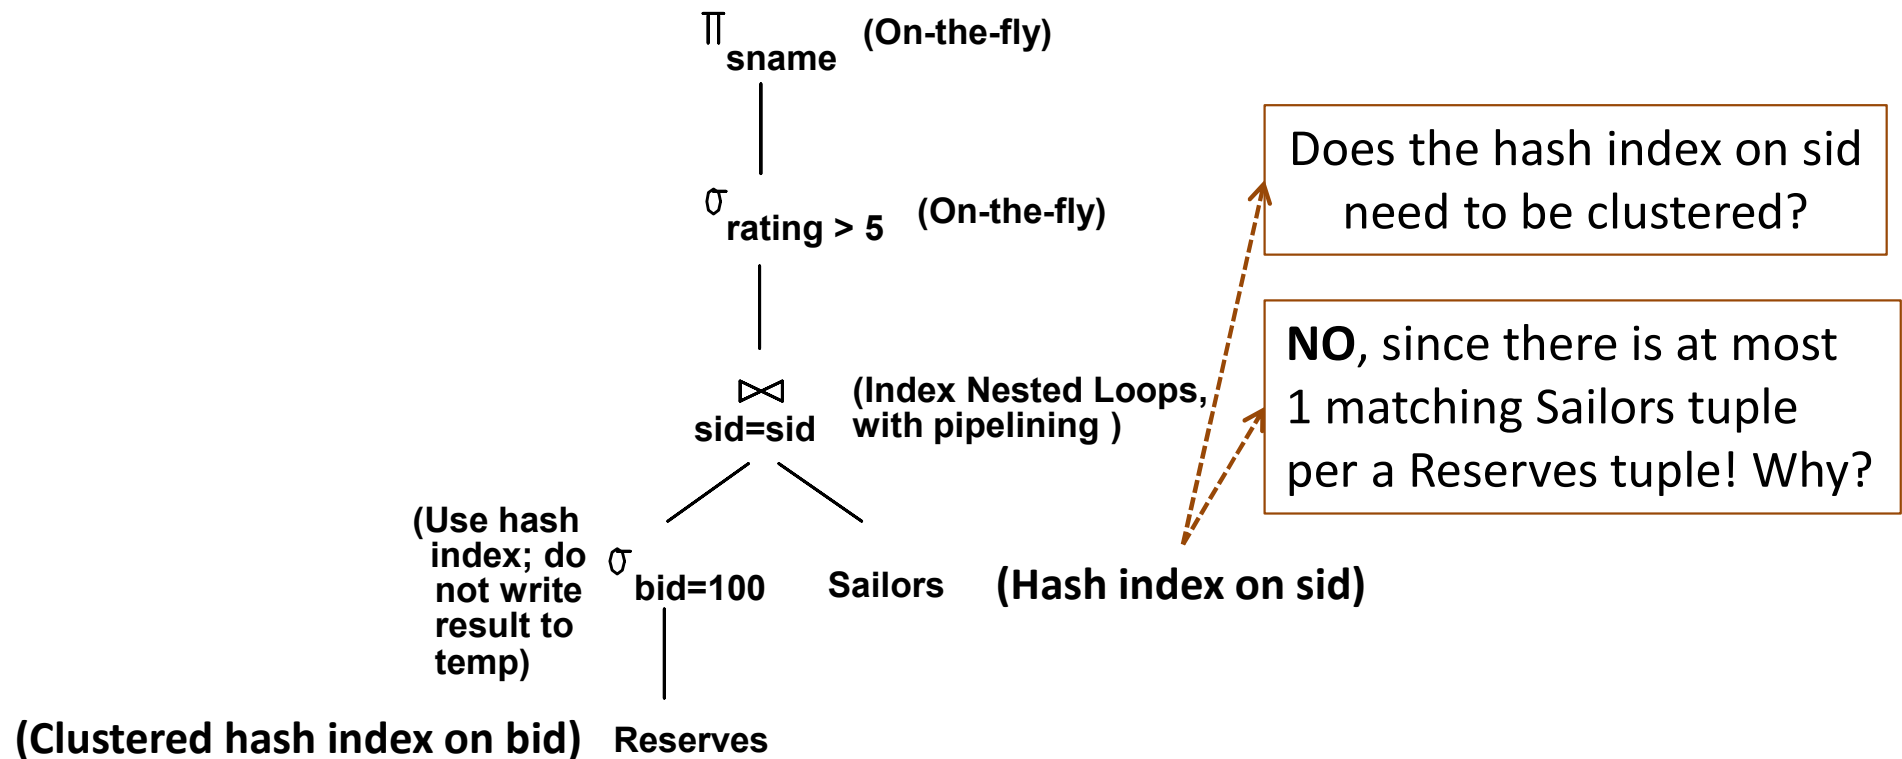
Is it necessary to project out unwanted columns?

**NO**, since selection results are NOT materialized



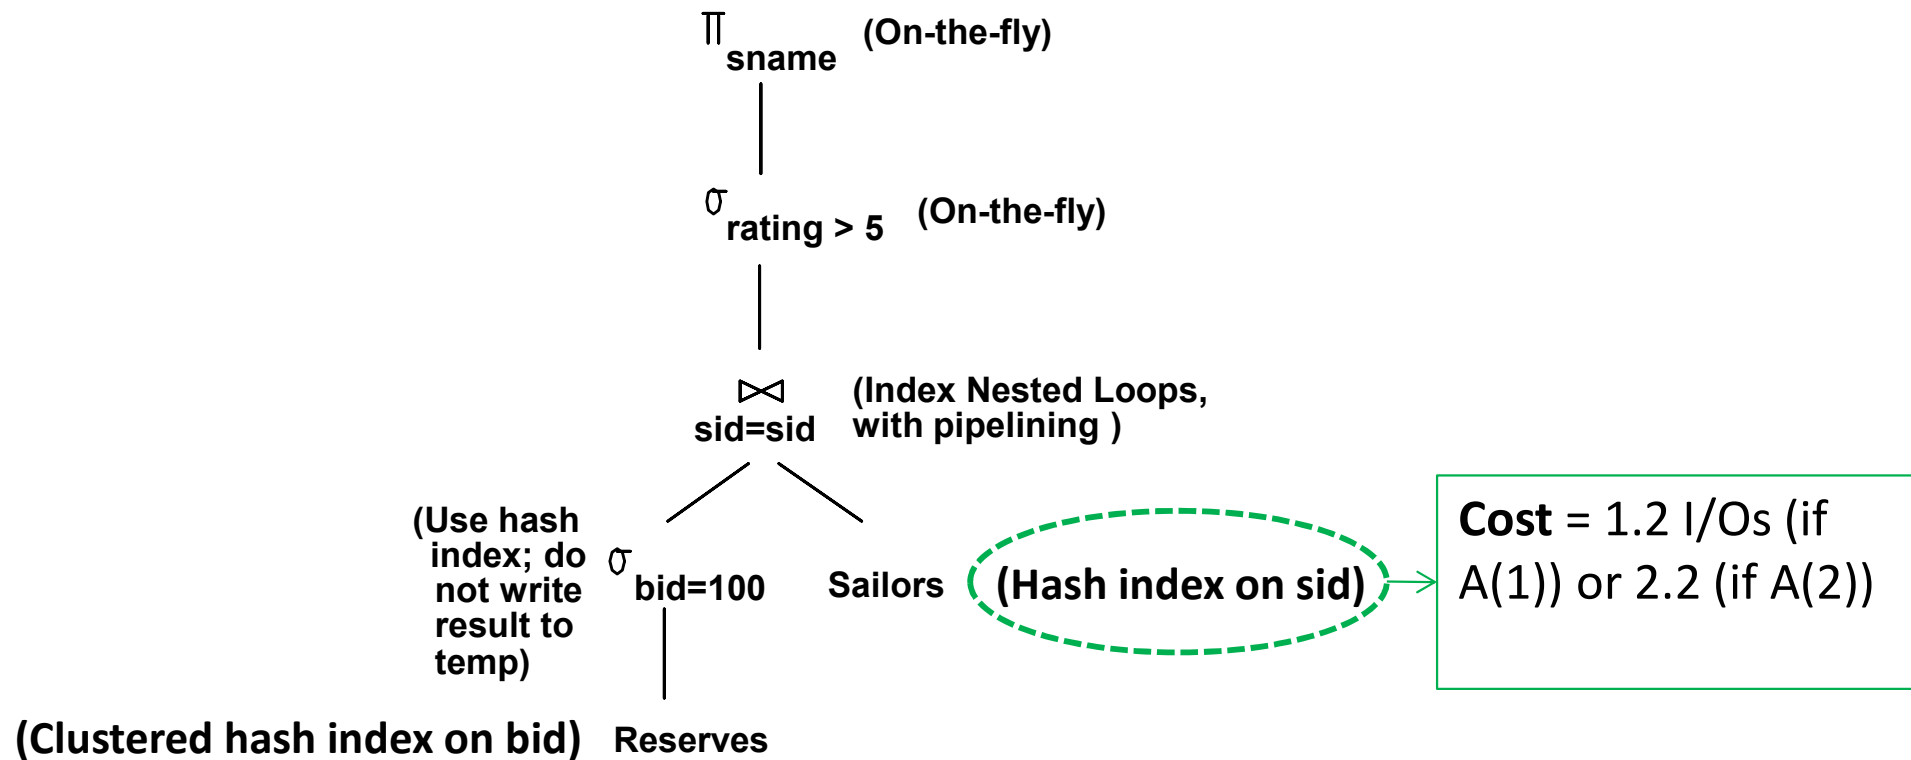
# Using Indexes

- What if indexes are available on Reserves and Sailors?



# Using Indexes

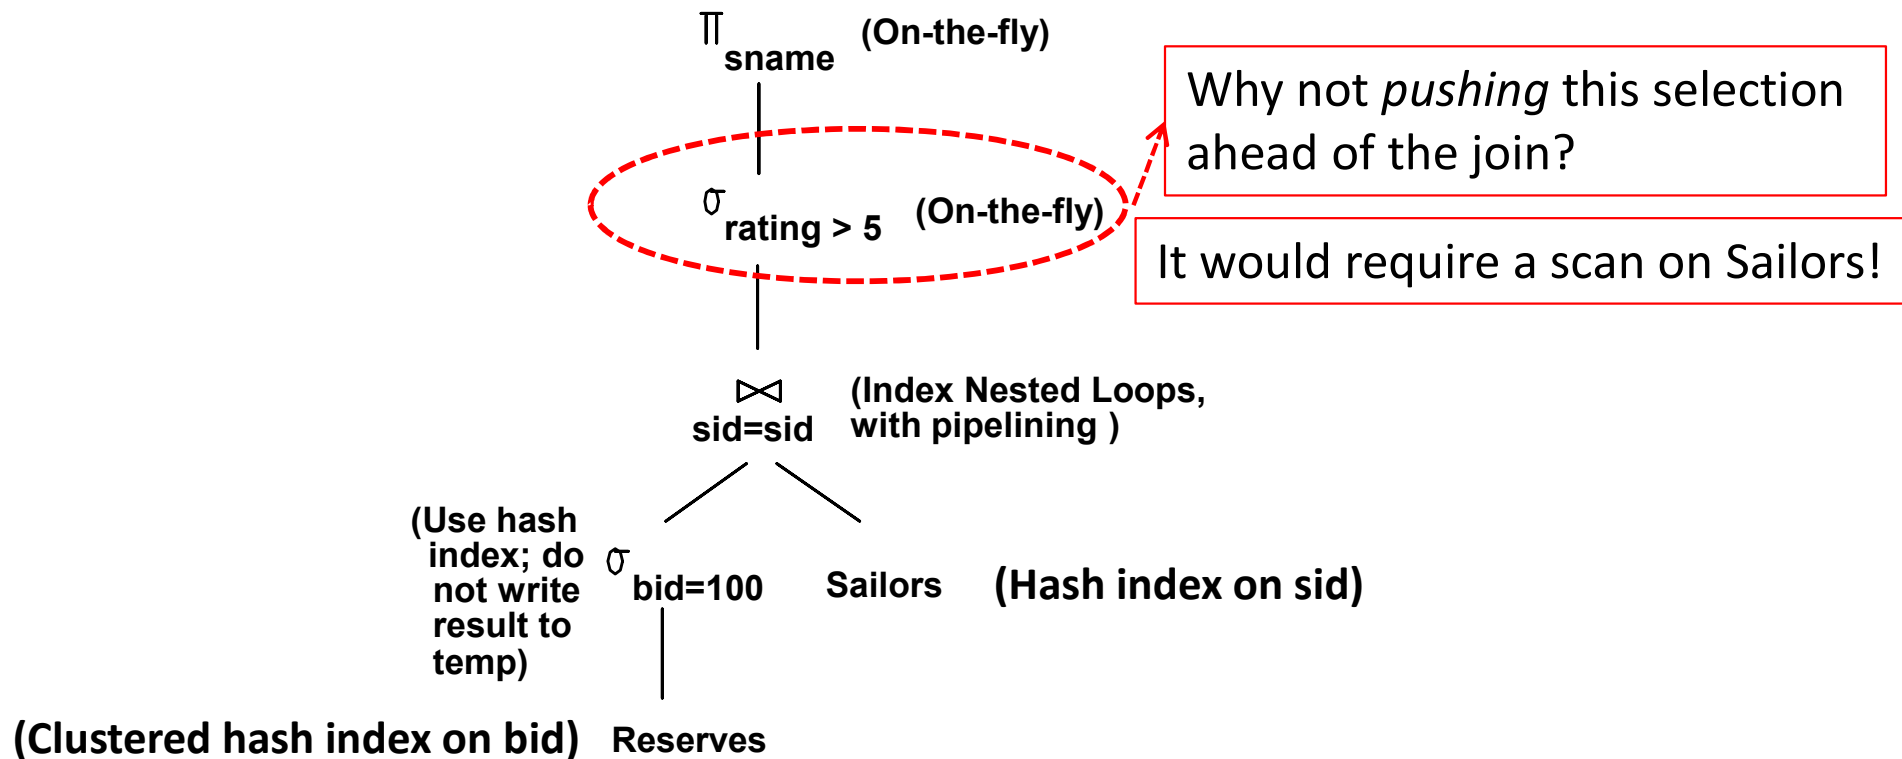
- What if indexes are available on Reserves and Sailors?





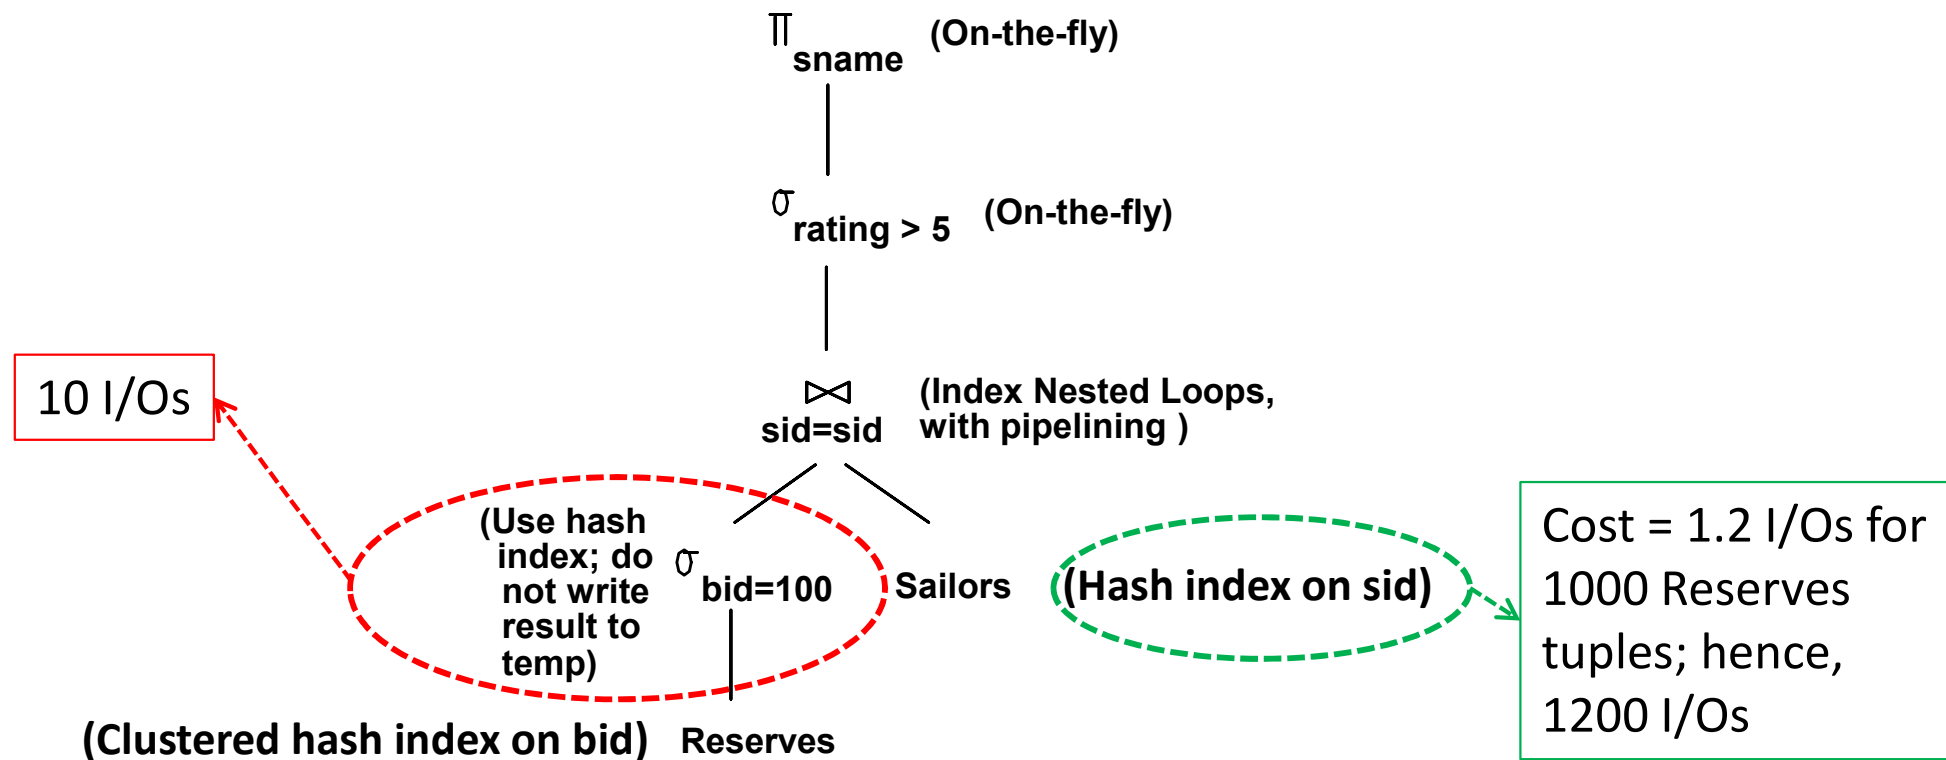
# Using Indexes

- What if indexes are available on Reserves and Sailors?



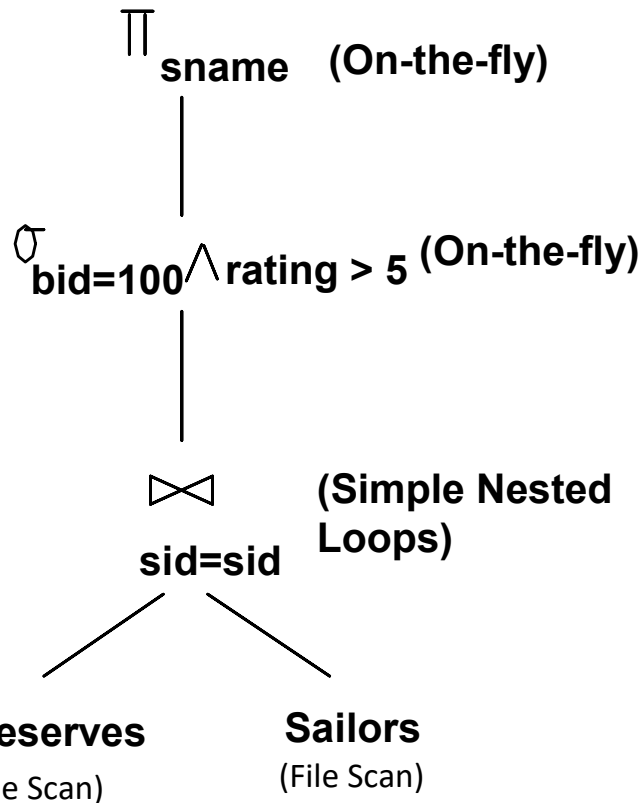
# The I/O Cost of the *New Q* Plan

- What is the I/O cost of the following evaluation plan?




**Total Cost = 10 + 1200 = 1210 I/Os**

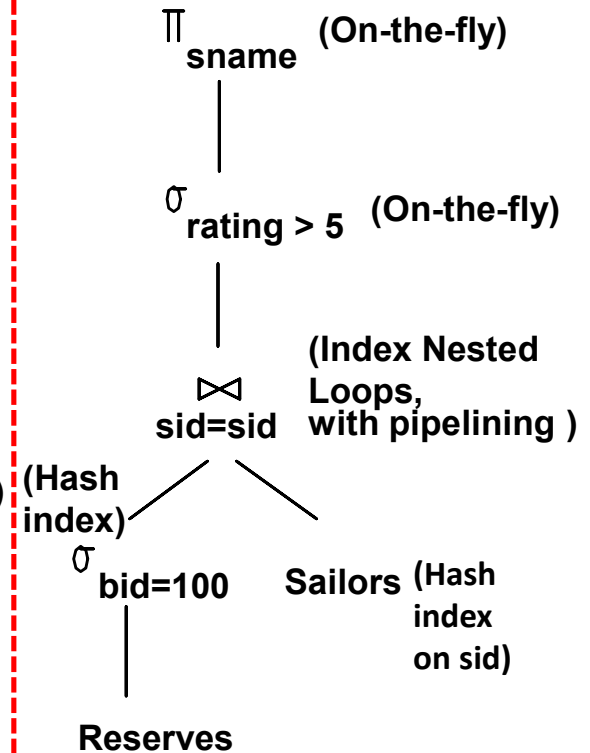
# Comparing I/O Costs: Recap



Total Cost = 501,000 I/Os

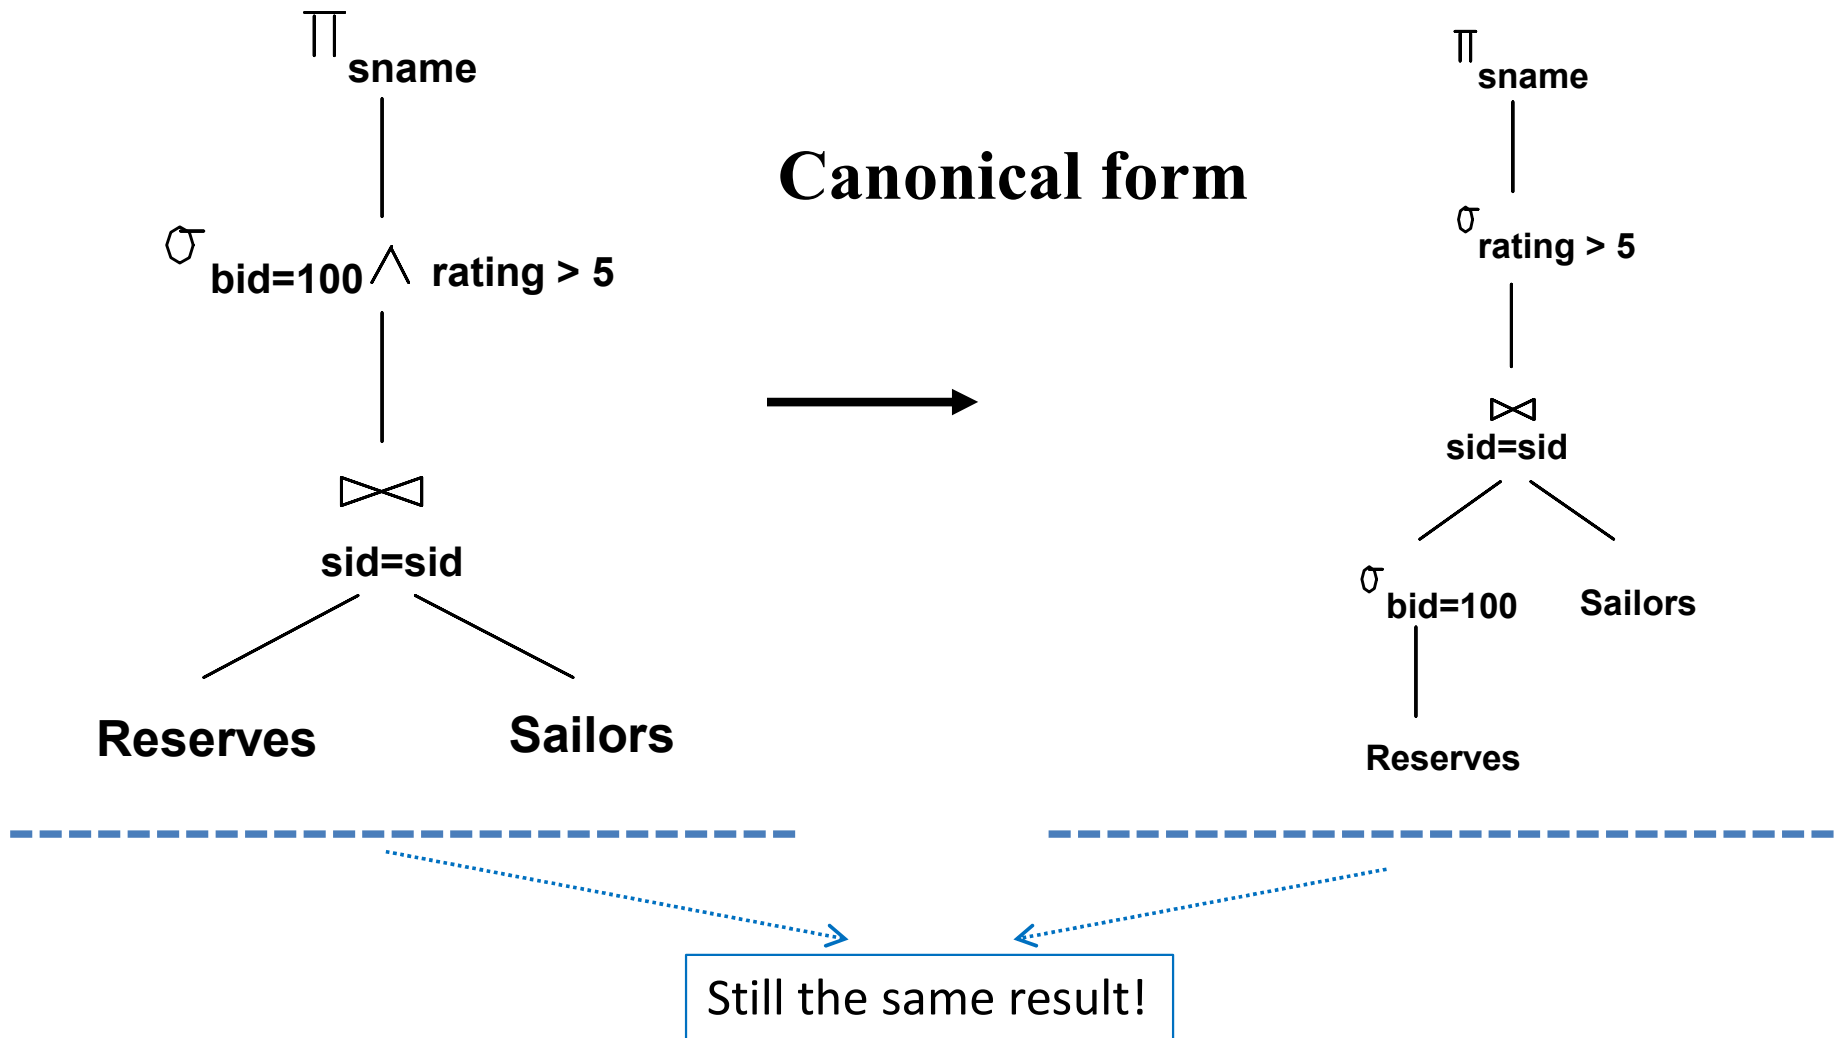


Total Cost = 4060 I/Os



Total Cost = 1210 I/Os

# But, How Can we Ensure Correctness?



How can this be guaranteed?

# Outline



A Brief Primer on Query Optimization

Evaluating Query Plans

Relational Algebra Equivalences

Estimating Plan Costs

Enumerating Plans



# Relational Algebra Equivalences

- A relational query optimizer uses *relational algebra equivalences* to identify many *equivalent* expressions for a given query
- Two relational algebra expressions over the same set of input relations are said to be *equivalent* if they produce the same result on all relations' instances
- Relational algebra equivalences allow us to:
  - Push selections and projections ahead of joins
  - Combine selections and cross-products into joins
  - Choose different join orders

# RA Equivalences: Selections

- Two important equivalences involve selections:

## 1. Cascading of Selections:

$$\sigma_{c1 \wedge \dots \wedge cn}(R) \equiv \sigma_{c1}(\dots \sigma_{cn}(R))$$

Allows us to combine several selections into one selection

**OR:** Allows us to replace a selection with several smaller selections

## 2. Commutation of Selections:

$$\sigma_{c1}(\sigma_{c2}(R)) \equiv \sigma_{c2}(\sigma_{c1}(R))$$

Allows us to test selection conditions in either order

# RA Equivalences: Projections

- One important equivalence involves projections:
  - Cascading of Projections:

$$\pi_{a_1}(R) \equiv \pi_{a_1}(\dots(\pi_{a_n}(R)))$$

This says that successively eliminating columns from a relation is equivalent to simply eliminating all but the columns retained by the final projection!



# RA Equivalences: Cross-Products and Joins

- Two important equivalences involve cross-products and joins:

## 1. Commutative Operations:

$$(R \times S) \equiv (S \times R)$$

$$(R \bowtie S) \equiv (S \bowtie R)$$

This allows us to choose which relation to be the inner and which to be the outer!

# RA Equivalences: Cross-Products and Joins

- Two important equivalences involve cross-products and joins:

## 2. Associative Operations:

$$R \times (S \times T) \equiv (R \times S) \times T$$

$$R \bowtie (S \bowtie T) \equiv (R \bowtie S) \bowtie T$$

It follows:  $R \bowtie (S \bowtie T) \equiv (T \bowtie R) \bowtie S$

This says that regardless of the order in which the relations are considered, the final result is the same!

This *order-independence* is fundamental to how a query optimizer generates alternative query evaluation plans

# RA Equivalences: Selections, Projections, Cross Products and Joins

- Selections with Projections:

$$\pi_a(\sigma_c(R)) \equiv \sigma_c(\pi_a(R))$$

This says we can commute a selection with a projection if the selection involves only attributes retained by the projection!

- Selections with Cross-Products:

$$R \bowtie_c T \equiv \sigma_c(R \times S)$$

This says we can combine a selection with a cross-product to form a join (*as per the definition of a join*)!

# RA Equivalences: Selections, Projections, Cross Products and Joins

- Selections with Cross-Products and with Joins:

$$\sigma_c(R \times S) \equiv \sigma_c(R) \times S$$

$$\sigma_c(R \bowtie S) \equiv \sigma_c(R) \bowtie S$$

**Caveat:** The attributes mentioned in  $c$  must appear only in  $R$  and *NOT* in  $S$

This says we can commute a selection with a cross-product or a join if the selection condition involves only attributes of one of the arguments to the cross-product or join!

# RA Equivalences: Selections, Projections, Cross Products and Joins

- Selections with Cross-Products and with Joins (*Cont'd*):

$$\begin{aligned}\sigma_c(R \times S) &\equiv \sigma_{c1 \wedge c2 \wedge c3}(R \times S) \\ &\equiv \sigma_{c1}(\sigma_{c2}(\sigma_{c3}(R \times S))) \\ &\equiv \sigma_{c1}(\sigma_{c2}(R) \times \sigma_{c3}(S))\end{aligned}$$

This says we can push part of the selection condition  $c$  ahead of the cross-product!

This applies to joins as well!

# RA Equivalences: Selections, Projections, Cross Products and Joins

- Projections with Cross-Products and with Joins:

$$\pi_a(R \times S) \equiv \pi_{a1}(R) \times \pi_{a2}(S)$$


$$\pi_a(R \bowtie_c S) \equiv \pi_{a1}(R) \bowtie_c \pi_{a2}(S)$$

$$\pi_a(R \bowtie_c S) \equiv \pi_a(\pi_{a1}(R) \bowtie_c \pi_{a2}(S))$$


Intuitively, we need to retain only those attributes of R and S that are either mentioned in the join condition **c** or included in the set of attributes **a** retained by the projection

# How to Estimate the Cost of Plans?

- Now that correctness is ensured, how can the DBMS estimate the costs of various plans?



Canonical form



# Next Class

